

Managing data consistency in a microservice architecture using Sagas

Chris Richardson

Founder of eventuate.io

Author of Microservices Patterns

Founder of the original CloudFoundry.com

Author of POJOs in Action

@crichardson

chris@chrisrichardson.net

<http://eventuate.io> <http://learn.microservices.io>

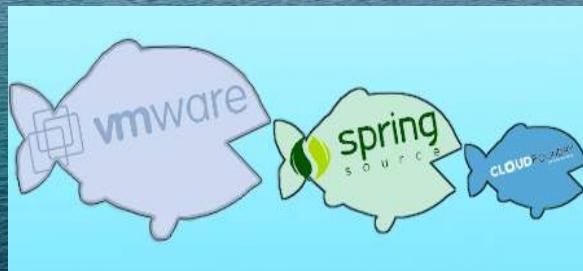
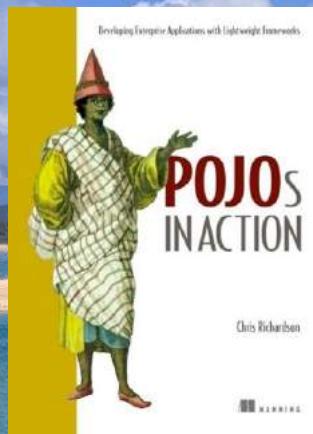
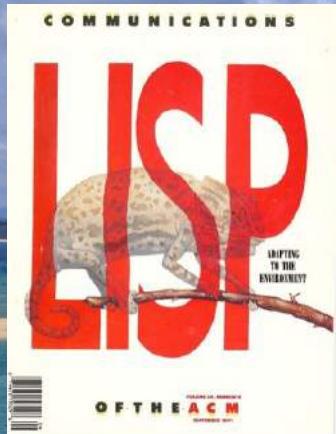
SATURN 2018

Presentation goal

Distributed data management challenges
in a microservice architecture

Sagas as the transaction model

About Chris



About Chris

Consultant and trainer
focusing on modern
application architectures
including microservices
(<http://www.chrisrichardson.net/>)

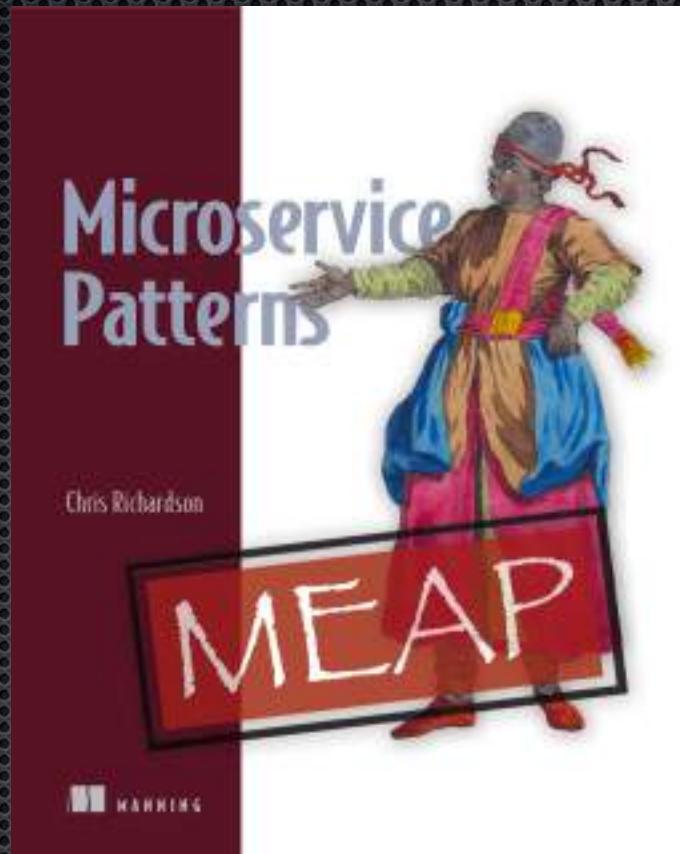
About Chris

Founder of a startup that is creating
an open-source/SaaS platform
that simplifies the development of
transactional microservices

(<http://eventuate.io>)



For more information



40%
discount
with code
ctwsaturn18

<http://learn.microservices.io>

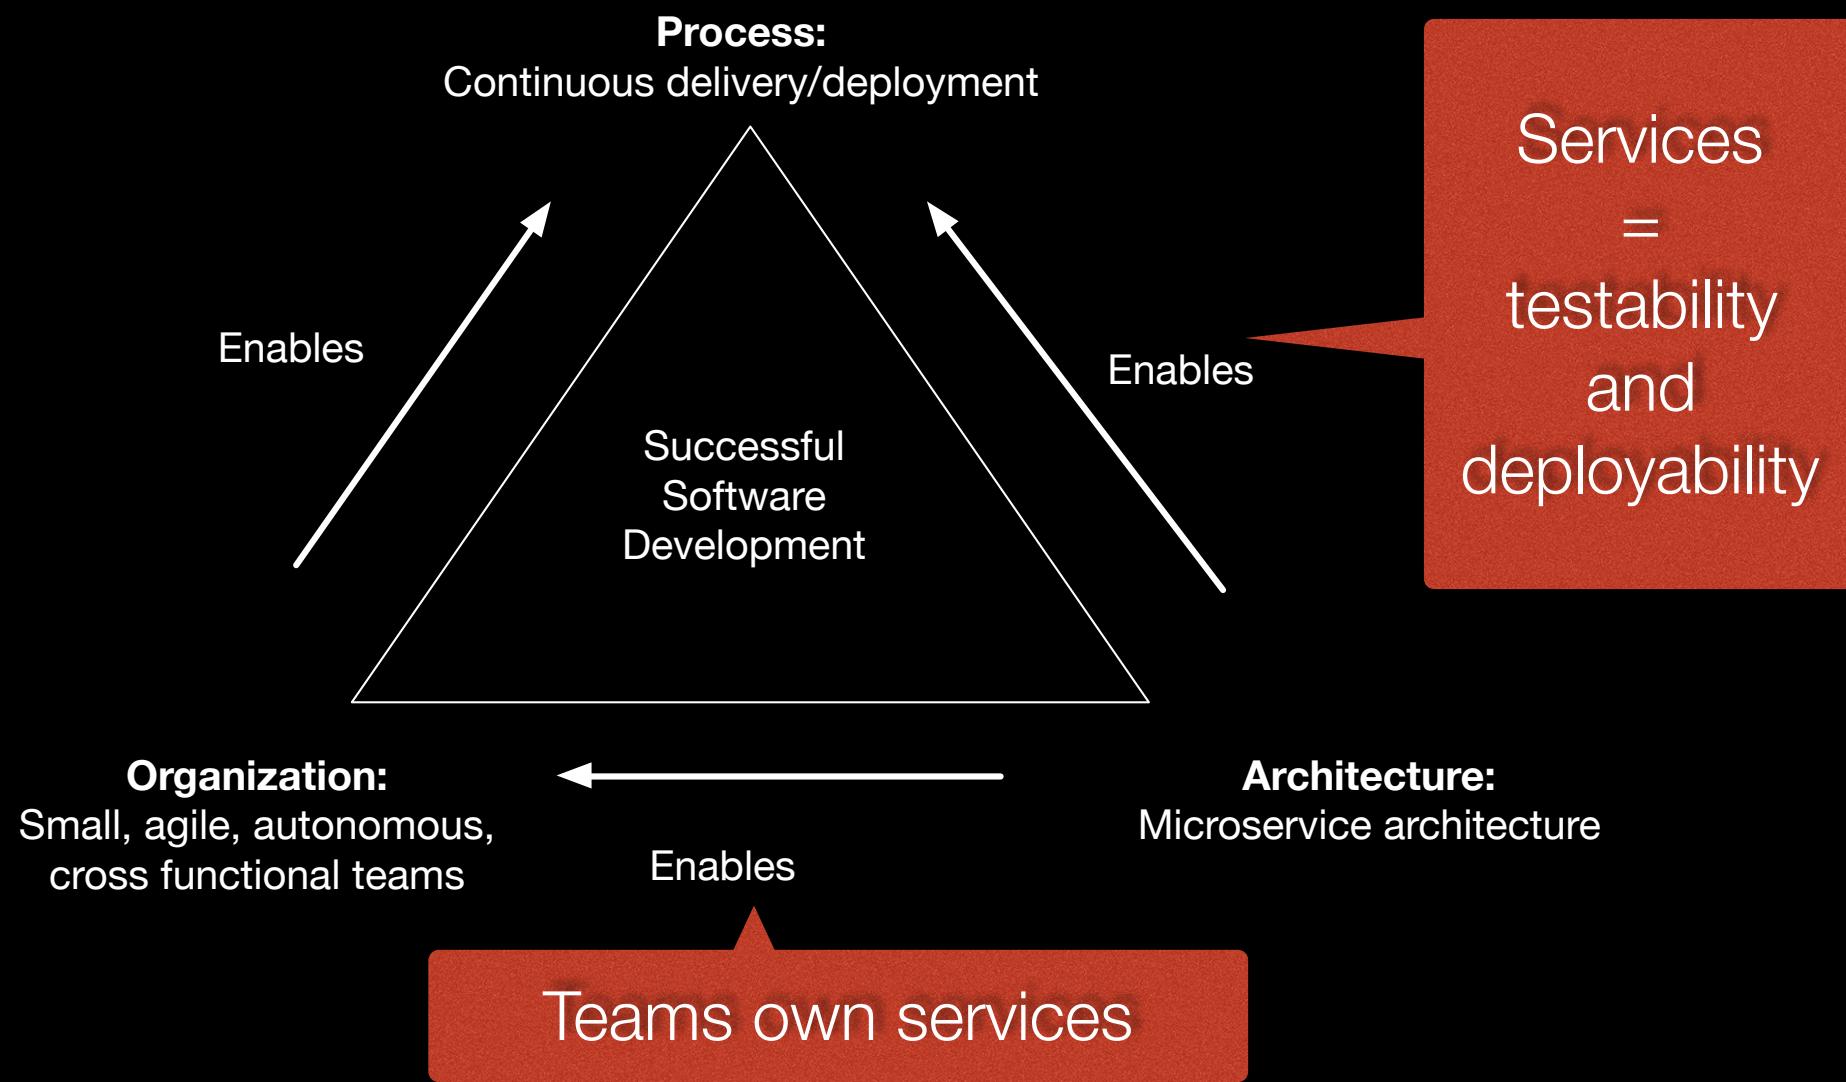
@crichtson

Agenda

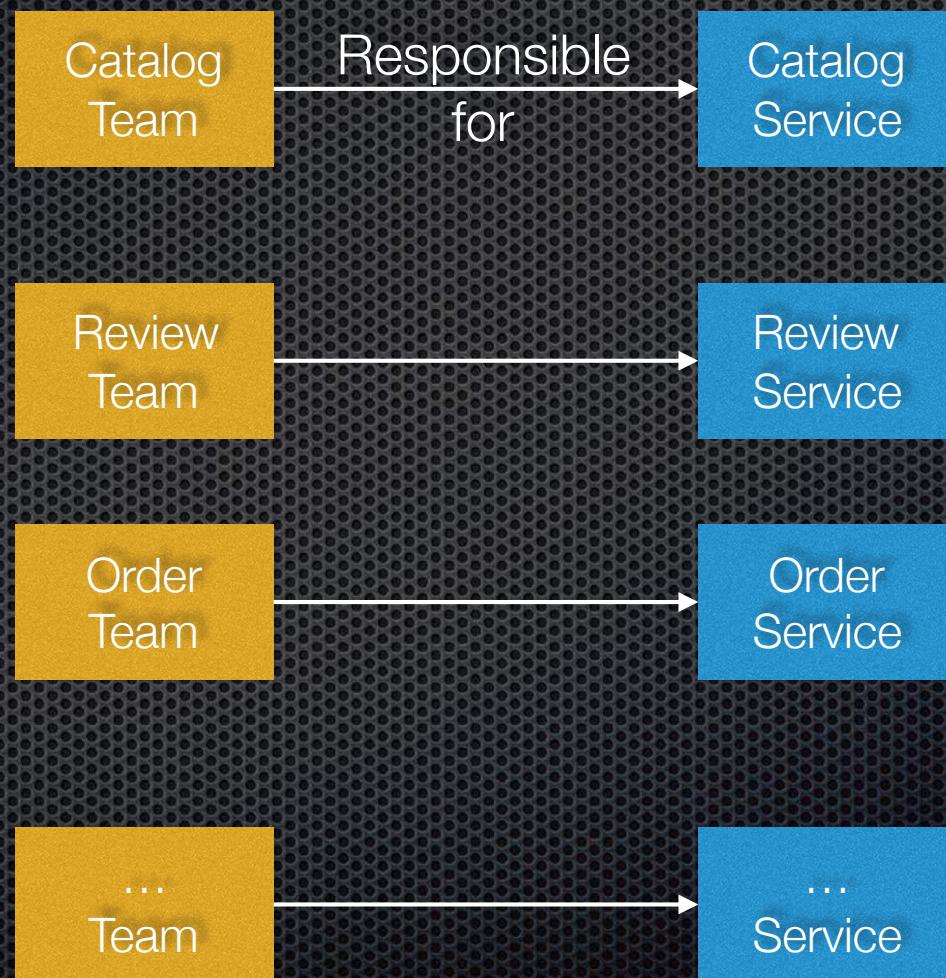
- ▣ ACID is not an option
- ▣ Overview of sagas
- ▣ Coordinating sagas
- ▣ Countermeasures for data anomalies
- ▣ Using reliable and transactional messaging

The microservice architecture
structures
an application as a
set of loosely coupled
services

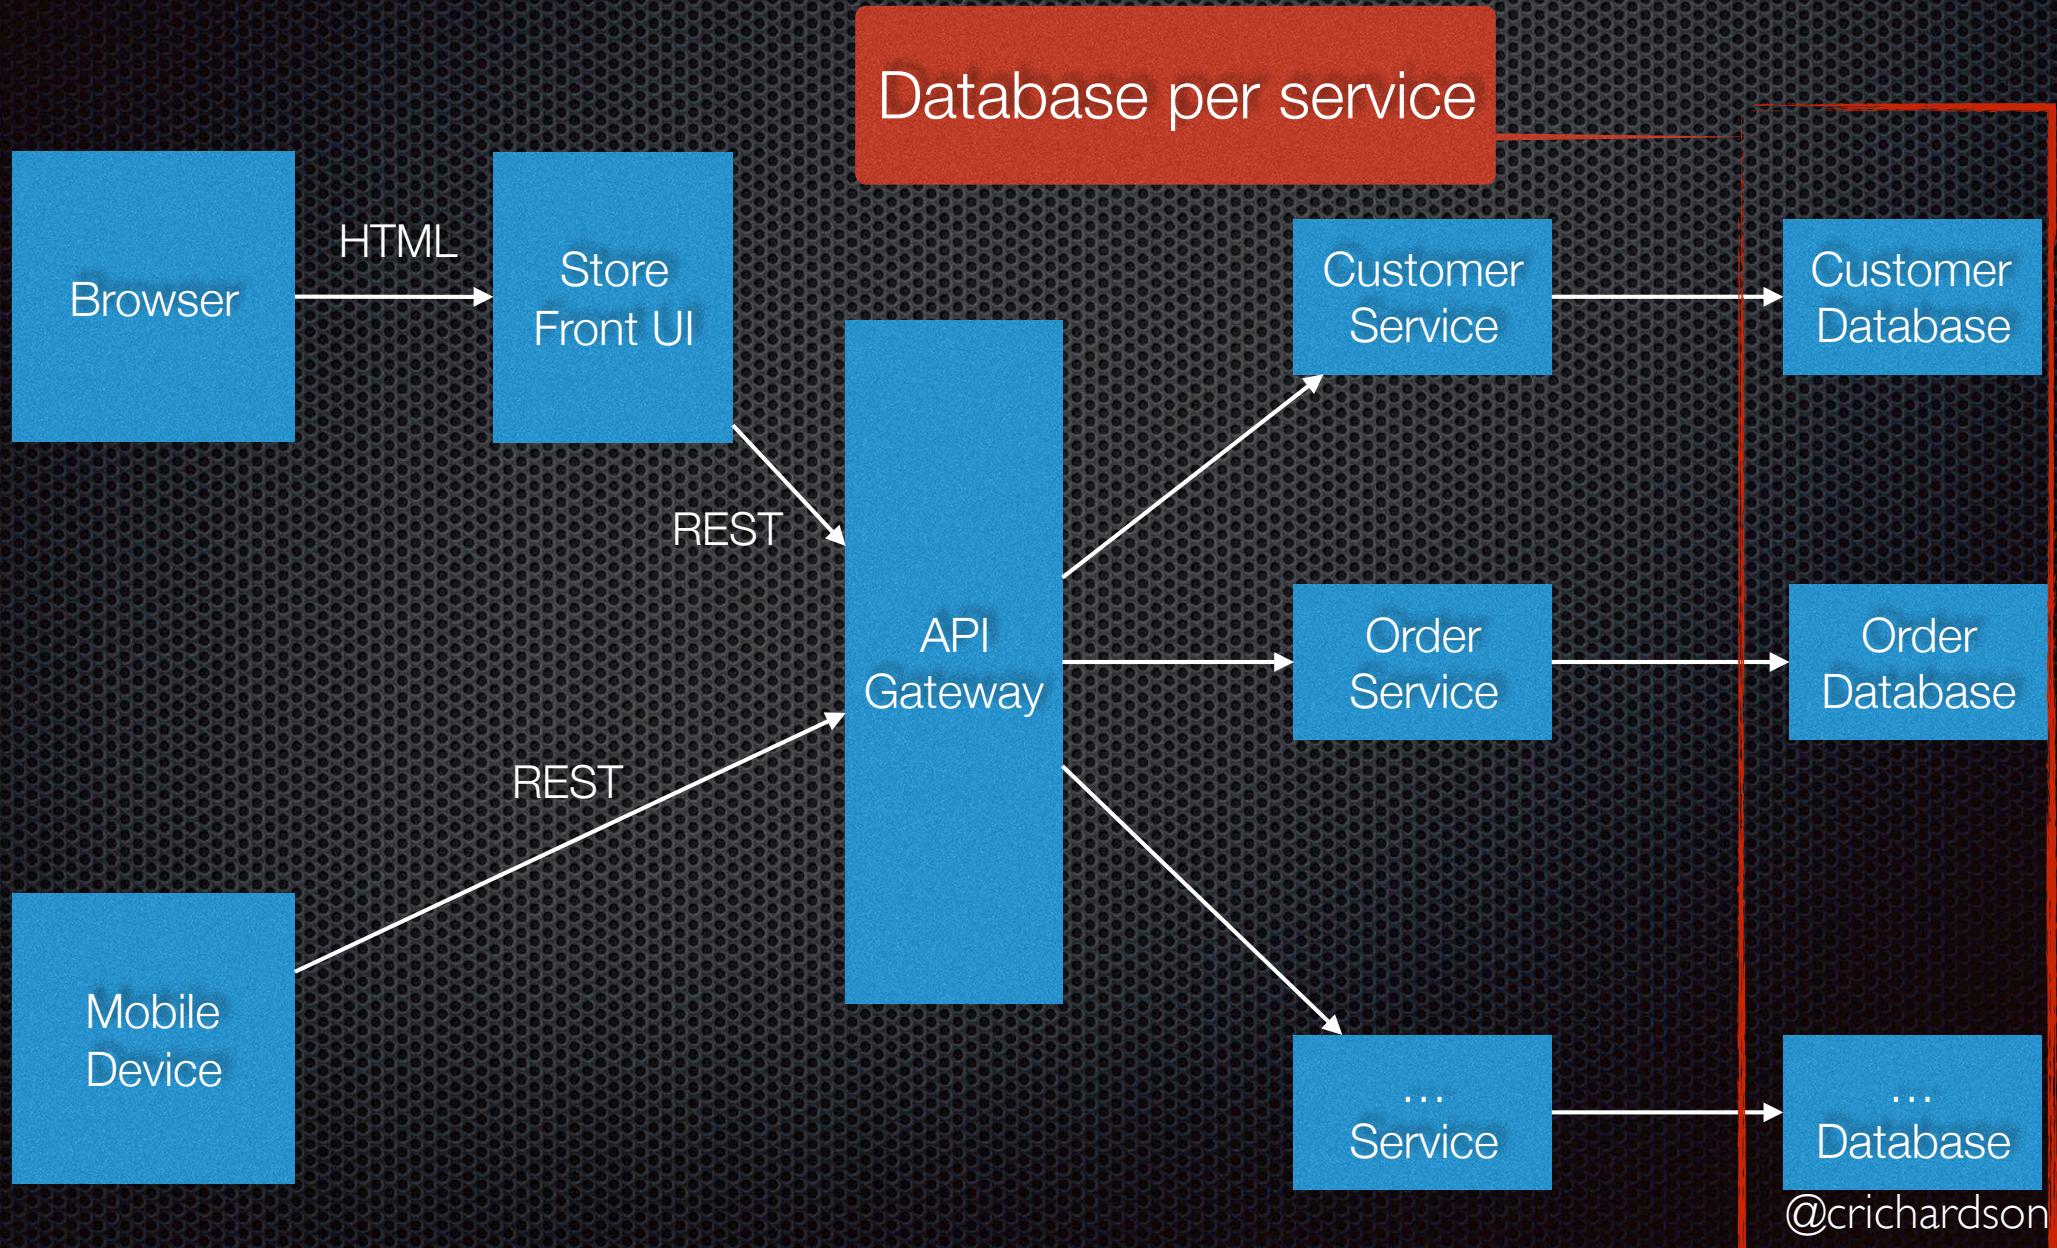
Microservices enable continuous delivery/deployment



Each team owns one or more services

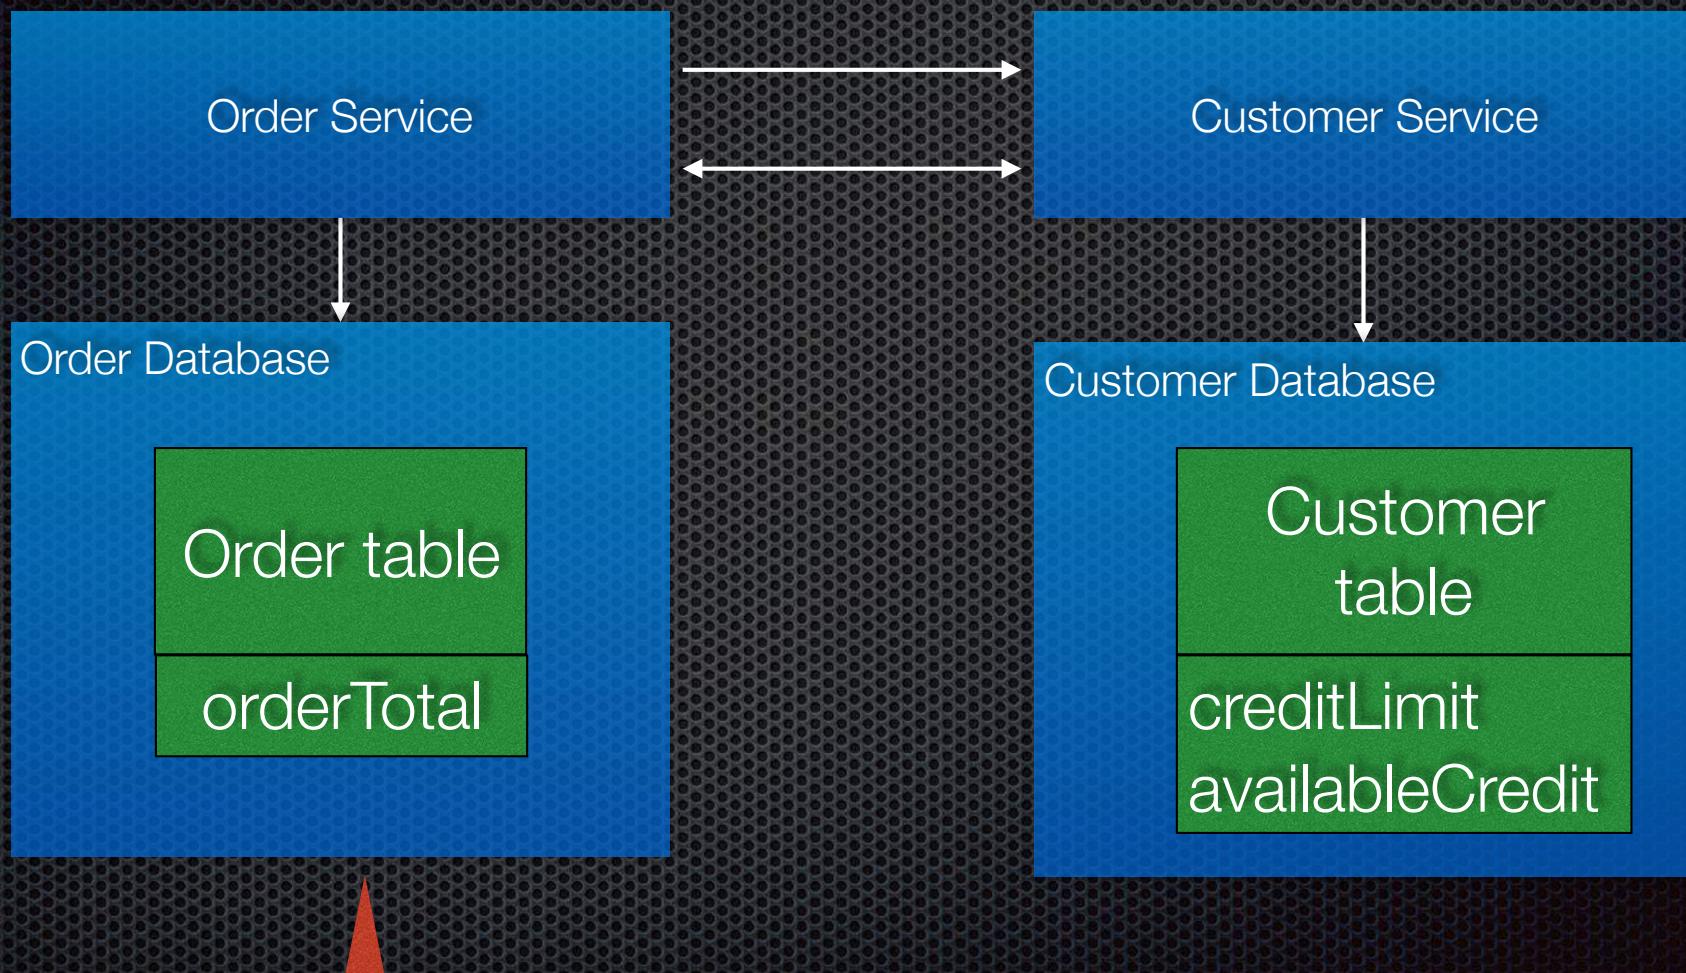


Microservice architecture



Private database
!=
private database **server**

Loose coupling = encapsulated data



Change schema without coordinating with other teams

But...

How to maintain data
consistency?

How to implement queries?

How to maintain data consistency?

```
createOrder(customerId, orderTotal)
```

Pre-conditions:

- customerId is valid

Post-conditions

- Order was created
- Customer.availableCredit -= orderTotal

Spans
services

Customer class

Invariant:

$\text{availableCredit} \geq 0$

$\text{availableCredit} \leq \text{creditLimit}$

Cannot use ACID transactions that span services

Distributed transactions

```
BEGIN TRANSACTION
```

```
...
```

```
SELECT ORDER_TOTAL  
FROM ORDERS WHERE CUSTOMER_ID = ?
```

```
...
```

```
SELECT CREDIT_LIMIT  
FROM CUSTOMERS WHERE CUSTOMER_ID = ?
```

```
...
```

```
INSERT INTO ORDERS ...
```

```
...
```

```
COMMIT TRANSACTION
```

Private to the
Order Service

Private to the
Customer Service

2PC is not an option

- Guarantees consistency

BUT

- 2PC coordinator is a single point of failure
- Chatty: at least $O(4n)$ messages, with retries $O(n^2)$
- Reduced throughput due to locks
- Not supported by many NoSQL databases (or message brokers)
- CAP theorem \Rightarrow 2PC impacts availability
-

ACID



Basically
Available
Soft state
Eventually consistent

Agenda

- ACID is not an option
- Overview of sagas
- Coordinating sagas
- Countermeasures for data anomalies
- Using reliable and transactional messaging

From a 1987 paper

SAGAS

*Hector Garcia-Molina
Kenneth Salem*

**Department of Computer Science
Princeton University
Princeton, N J 08544**

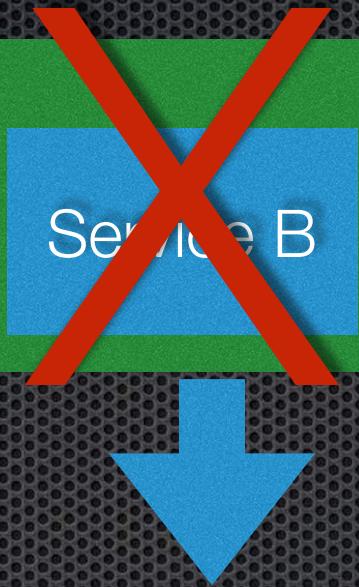
Use Sagas instead of 2PC

Distributed transaction

Service A

Service B

Service C



Saga

Service A

Local
transaction

Service B

Local
transaction

Service C

Local
transaction

Create Order Saga

createOrder()

Initiates saga

Order Service

Local transaction

createOrder()

Order

state=PENDING

Customer Service

Local transaction

reserveCredit()

Customer

Order Service

Local transaction

approve
order()

Order

state=APPROVED

But what about rollback?

BEGIN TRANSACTION

...
UPDATE ...

...
INSERT

.... BUSINESS RULE VIOLATED!!!!

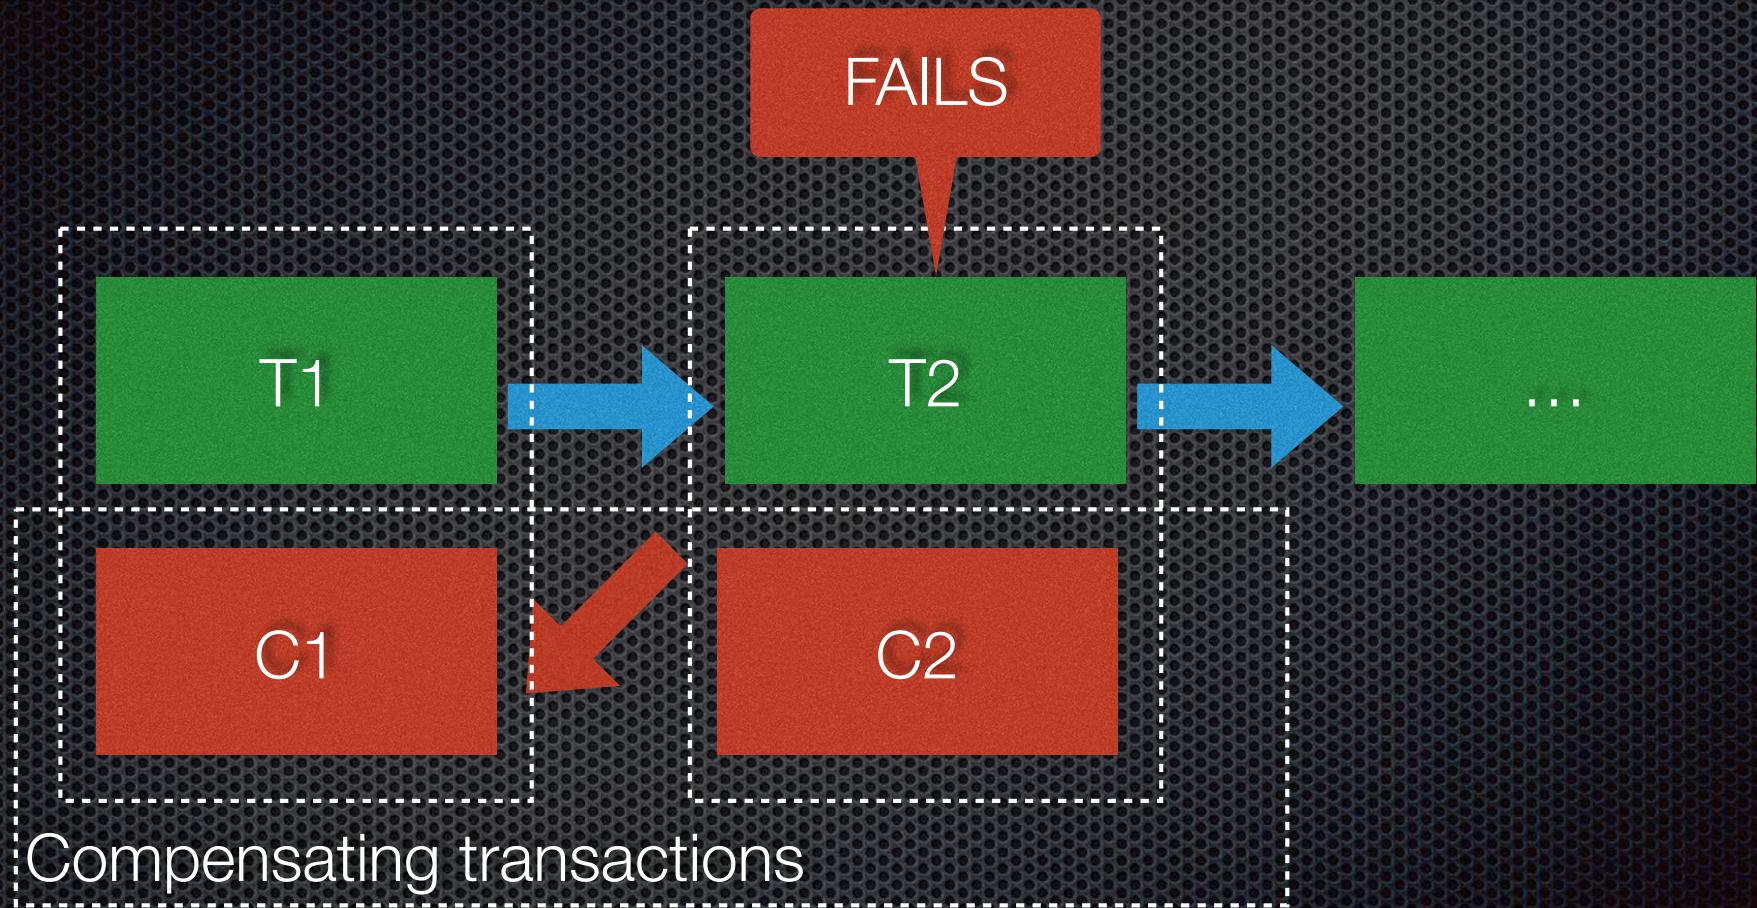
...
ROLLBACK TRANSACTION

Really simple!

Rolling back sagas

- ▣ Use compensating transactions
- ▣ Developer must write application logic to “rollback” eventually consistent transactions
- ▣ Careful design required!

Saga: Every Ti has a Ci



$T1 \Rightarrow T2 \Rightarrow C1$

Create Order Saga - rollback

createOrder()



Order Service

Local transaction

createOrder()



Order



Customer Service

Local transaction

reserveCredit()



Customer

FAIL



Order Service

Local transaction

reject
order()

Order

Insufficient credit

Writing compensating transactions isn't always easy

- Write them so they will always succeed
 - If a compensating transaction fails => no clear way to recover
- Challenge: Undoing changes when data has already been changed by a different transaction/saga
 - More on this later

Non-compensatable actions

- For example, sending an email can't be unsent.
- Move actions that can't be undone to the end of the saga
- More on this later.

Sagas complicate API design

- ✖ Synchronous API vs Asynchronous Saga
- ✖ Request initiates the saga. When to send back the response?
- ✖ Option #1: Send response when saga completes:
 - + Response specifies the outcome
 - Reduced availability
- ✖ Option #2: Send response immediately after creating the saga
(recommended):
 - + Improved availability
 - Response does not specify the outcome. Client must poll or be notified

Revised Create Order API

- ▣ createOrder()
 - ▣ returns id of newly created order
 - ▣ **NOT** fully validated
- ▣ getOrder(id)
 - ▣ Called periodically by client to get outcome of validation

Minimal impact on UI

- UI hides asynchronous API from the user
- Saga will usually appear instantaneous ($\leq 100\text{ms}$)
- If it takes longer \Rightarrow UI displays “processing” popup
- Server can push notification to UI

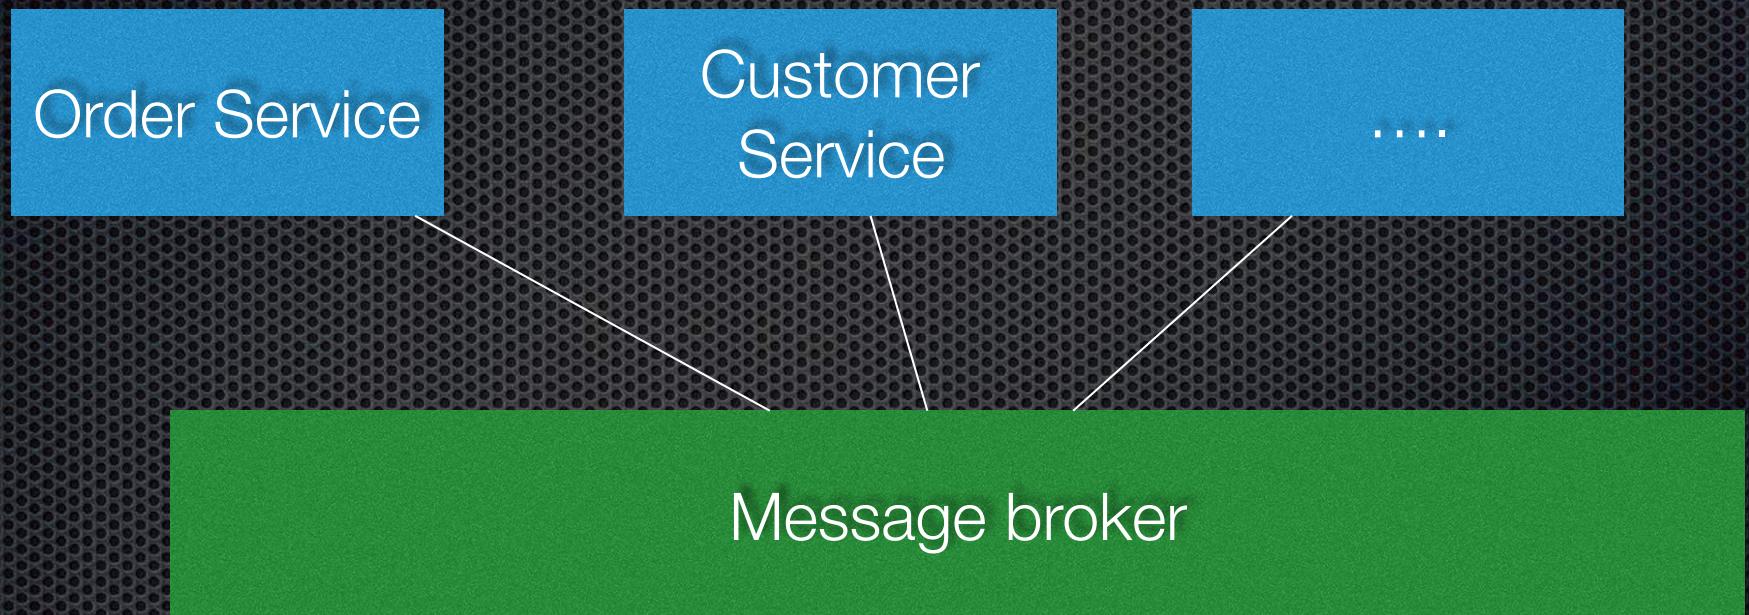
Agenda

- ▣ ACID is not an option
- ▣ Overview of sagas
- ▣ Coordinating sagas
- ▣ Countermeasures for data anomalies
- ▣ Using reliable and transactional messaging

How to sequence the saga transactions?

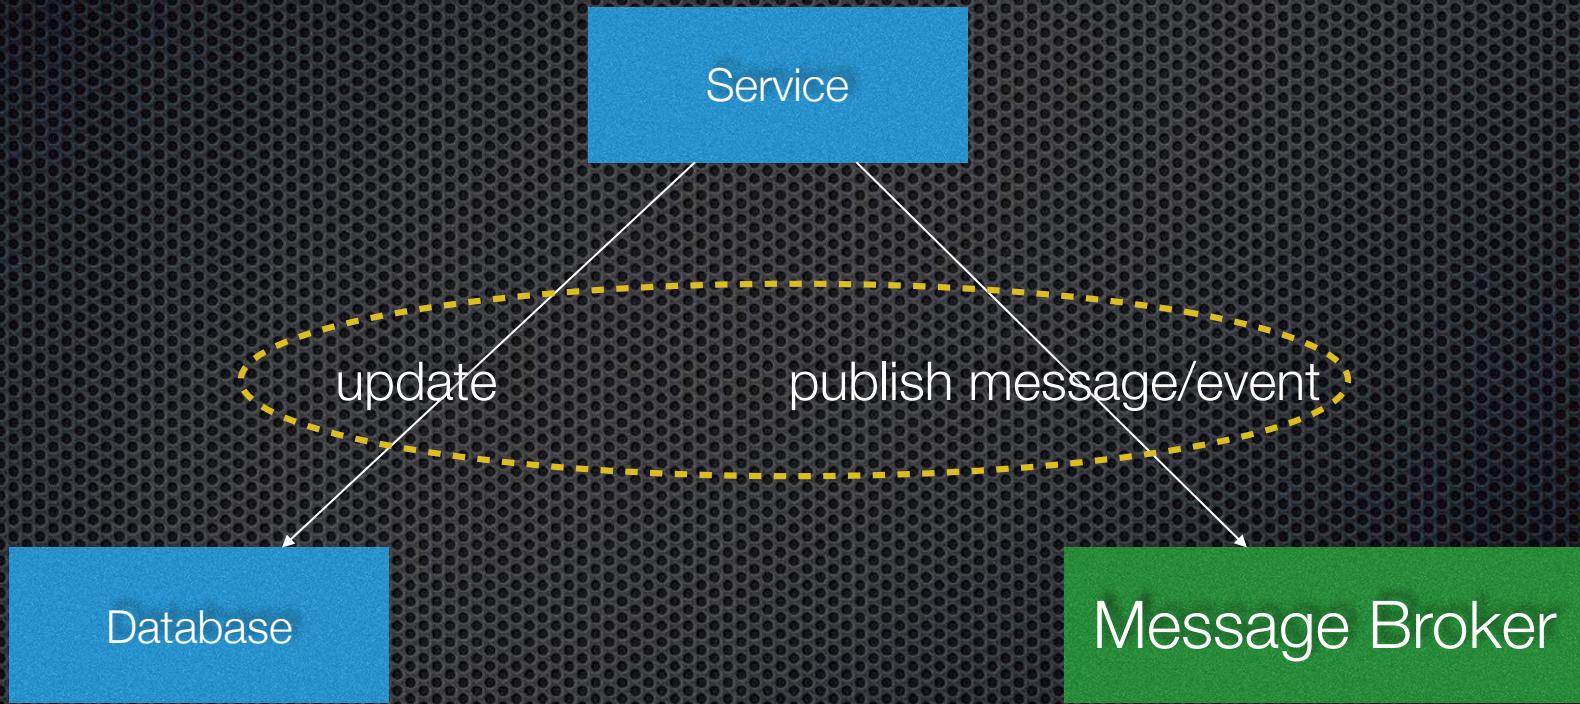
- After the completion of transaction T_i “something” must decide what step to execute next
- Success: which $T(i+1)$ - branching
- Failure: $C(i - 1)$

Use asynchronous, broker-based messaging



Guaranteed delivery ensures a saga complete when its participants are temporarily unavailable

Saga step = a transaction local to a service



Transactional DB: BEGIN ... COMMIT

DDD: Aggregate

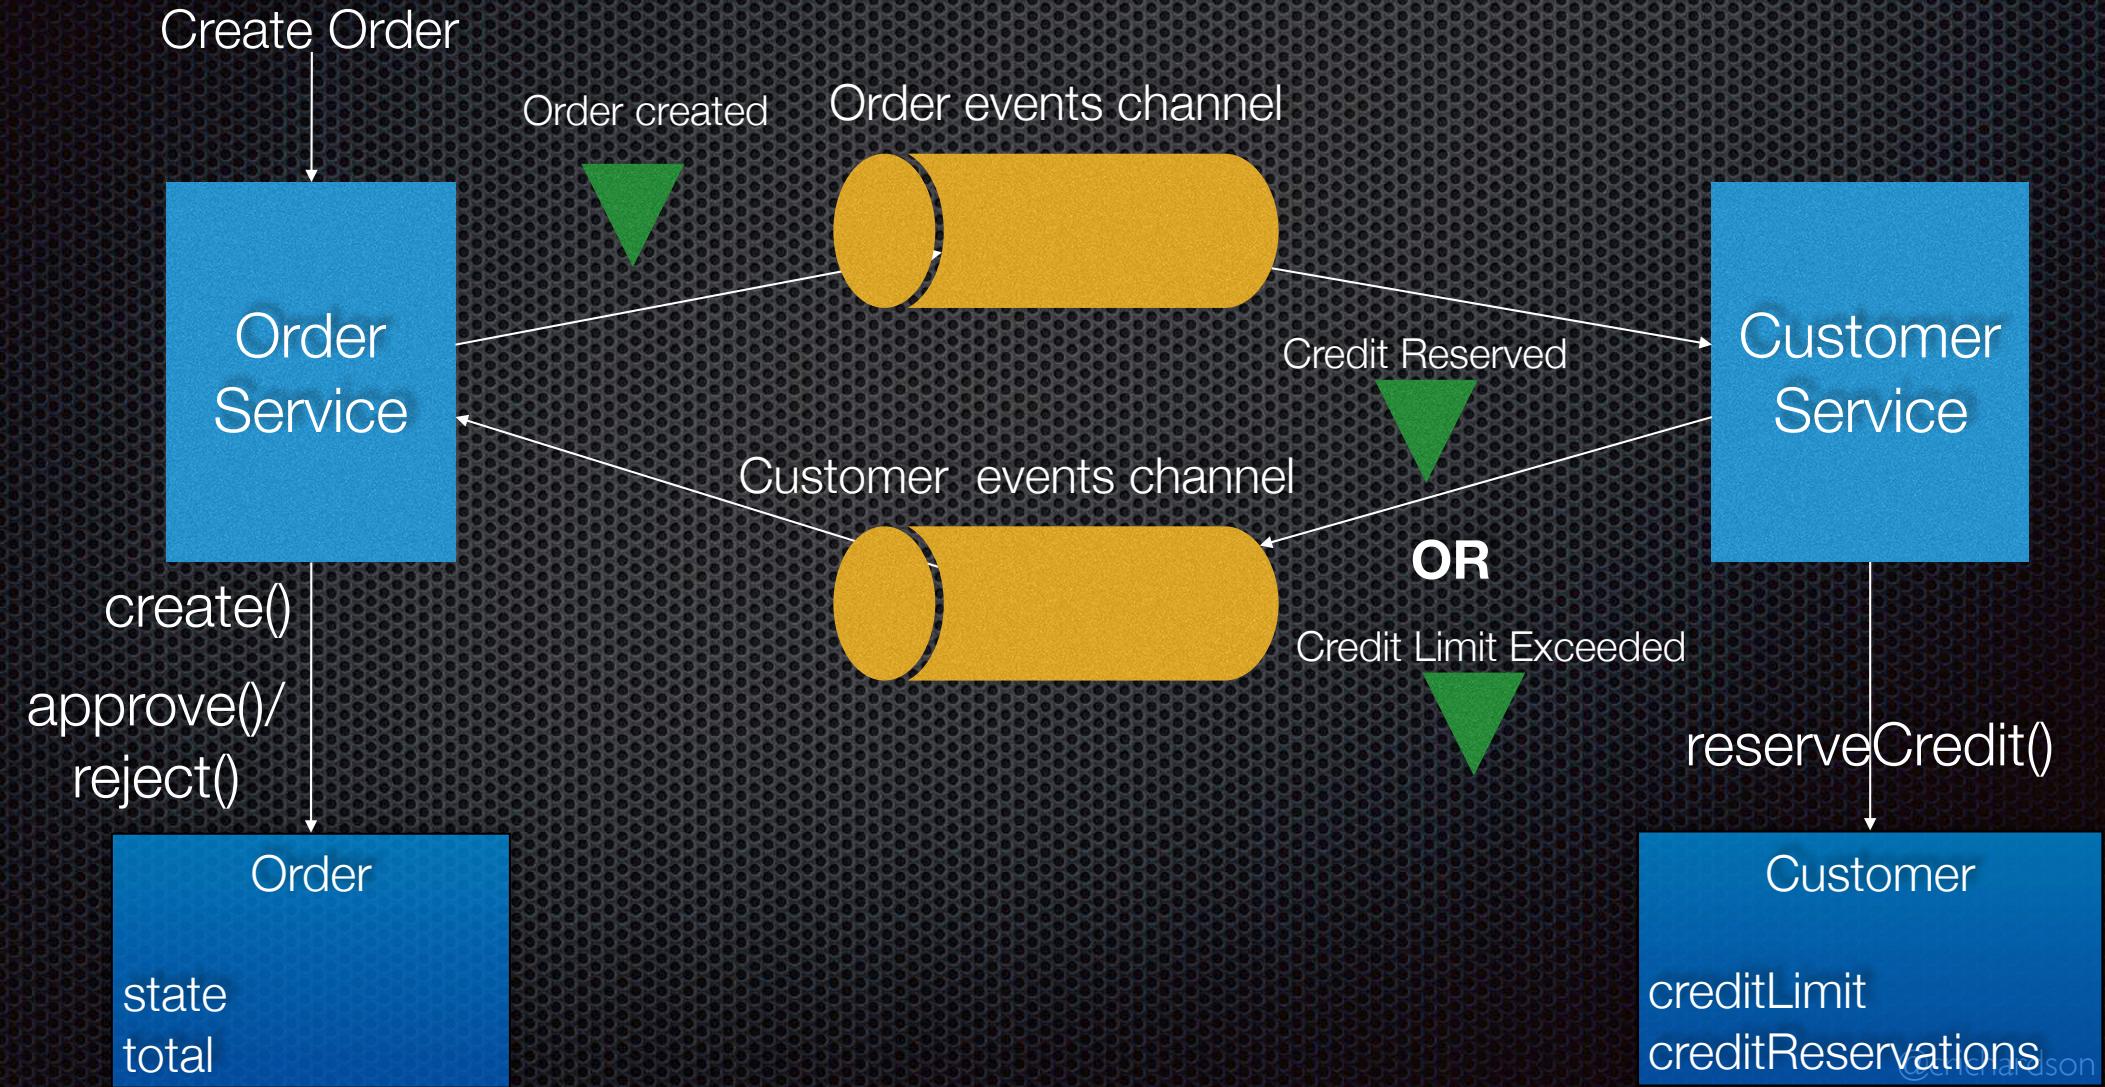
NoSQL: single 'record'

Choreography: **distributed** decision making

vs.

Orchestration: **centralized** decision making

Option #1: Choreography-based coordination using events



Order Service: publishing domain events

```
public class OrderService {  
  
    @Autowired  
    private DomainEventPublisher domainEventPublisher;  
  
    @Autowired  
    private OrderRepository orderRepository;  
  
    @Transactional  
    public Order createOrder(OrderDetails orderDetails) {  
        Order order = Order.createOrder(orderDetails);  
        orderRepository.save(order);  
        domainEventPublisher.publish(Order.class,  
            order.getId(),  
            singletonList(new OrderCreatedEvent(order.getId(), orderDetails)));  
        return order;  
    }  
}
```

Create order

Publish event

Customer Service: consuming domain events...

```
public class OrderEventConsumer {

    @Autowired
    private CustomerRepository customerRepository;

    @Autowired
    private DomainEventPublisher domainEventPublisher;

    public DomainEventHandlers domainEventHandlers() {
        return DomainEventHandlersBuilder
            .forAggregateType("io.eventuate.examples.tram.ordersandcustomers.orders.domain.Order")
            .onEvent(OrderCreatedEvent.class, this::orderCreatedEventHandler)
            .build();
    }

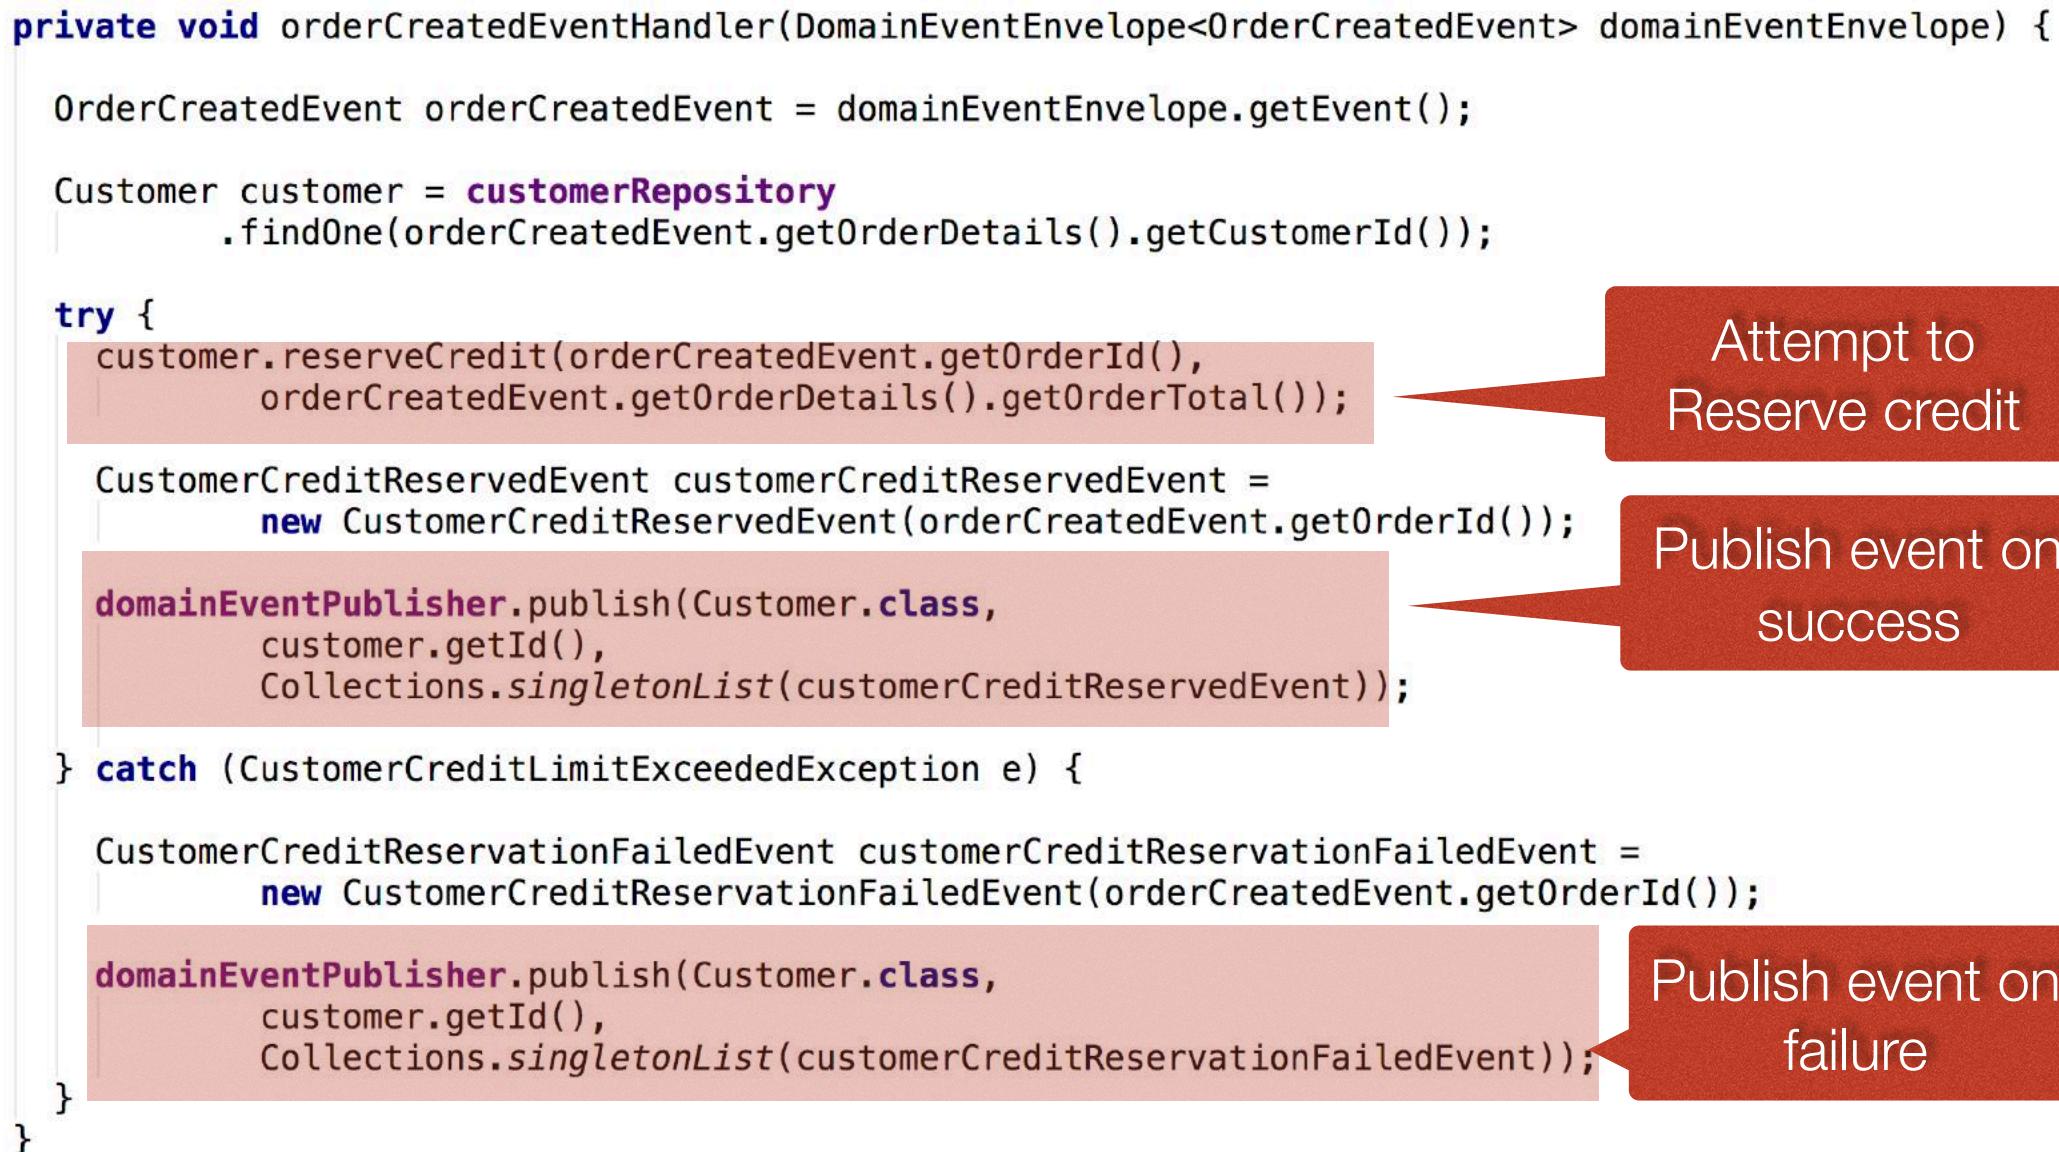
    private void orderCreatedEventHandler(DomainEventEnvelope<OrderCreatedEvent> domainEventEnvelope) {...}
}
```

Subscribe to event

<https://github.com/eventuate-tram/eventuate-tram-examples-customers-and-orders>

Customer Service: consuming domain events

```
private void orderCreatedEventHandler(DomainEventEnvelope<OrderCreatedEvent> domainEventEnvelope) {  
  
    OrderCreatedEvent orderCreatedEvent = domainEventEnvelope.getEvent();  
  
    Customer customer = customerRepository  
        .findOne(orderCreatedEvent.getOrderDetails().getCustomerId());  
  
    try {  
        customer.reserveCredit(orderCreatedEvent.getOrderId(),  
            orderCreatedEvent.getOrderDetails().getOrderTotal());  
  
        CustomerCreditReservedEvent customerCreditReservedEvent =  
            new CustomerCreditReservedEvent(orderCreatedEvent.getOrderId());  
  
        domainEventPublisher.publish(Customer.class,  
            customer.getId(),  
            Collections.singletonList(customerCreditReservedEvent));  
  
    } catch (CustomerCreditLimitExceededException e) {  
  
        CustomerCreditReservationFailedEvent customerCreditReservationFailedEvent =  
            new CustomerCreditReservationFailedEvent(orderCreatedEvent.getOrderId());  
  
        domainEventPublisher.publish(Customer.class,  
            customer.getId(),  
            Collections.singletonList(customerCreditReservationFailedEvent));  
    }  
}
```

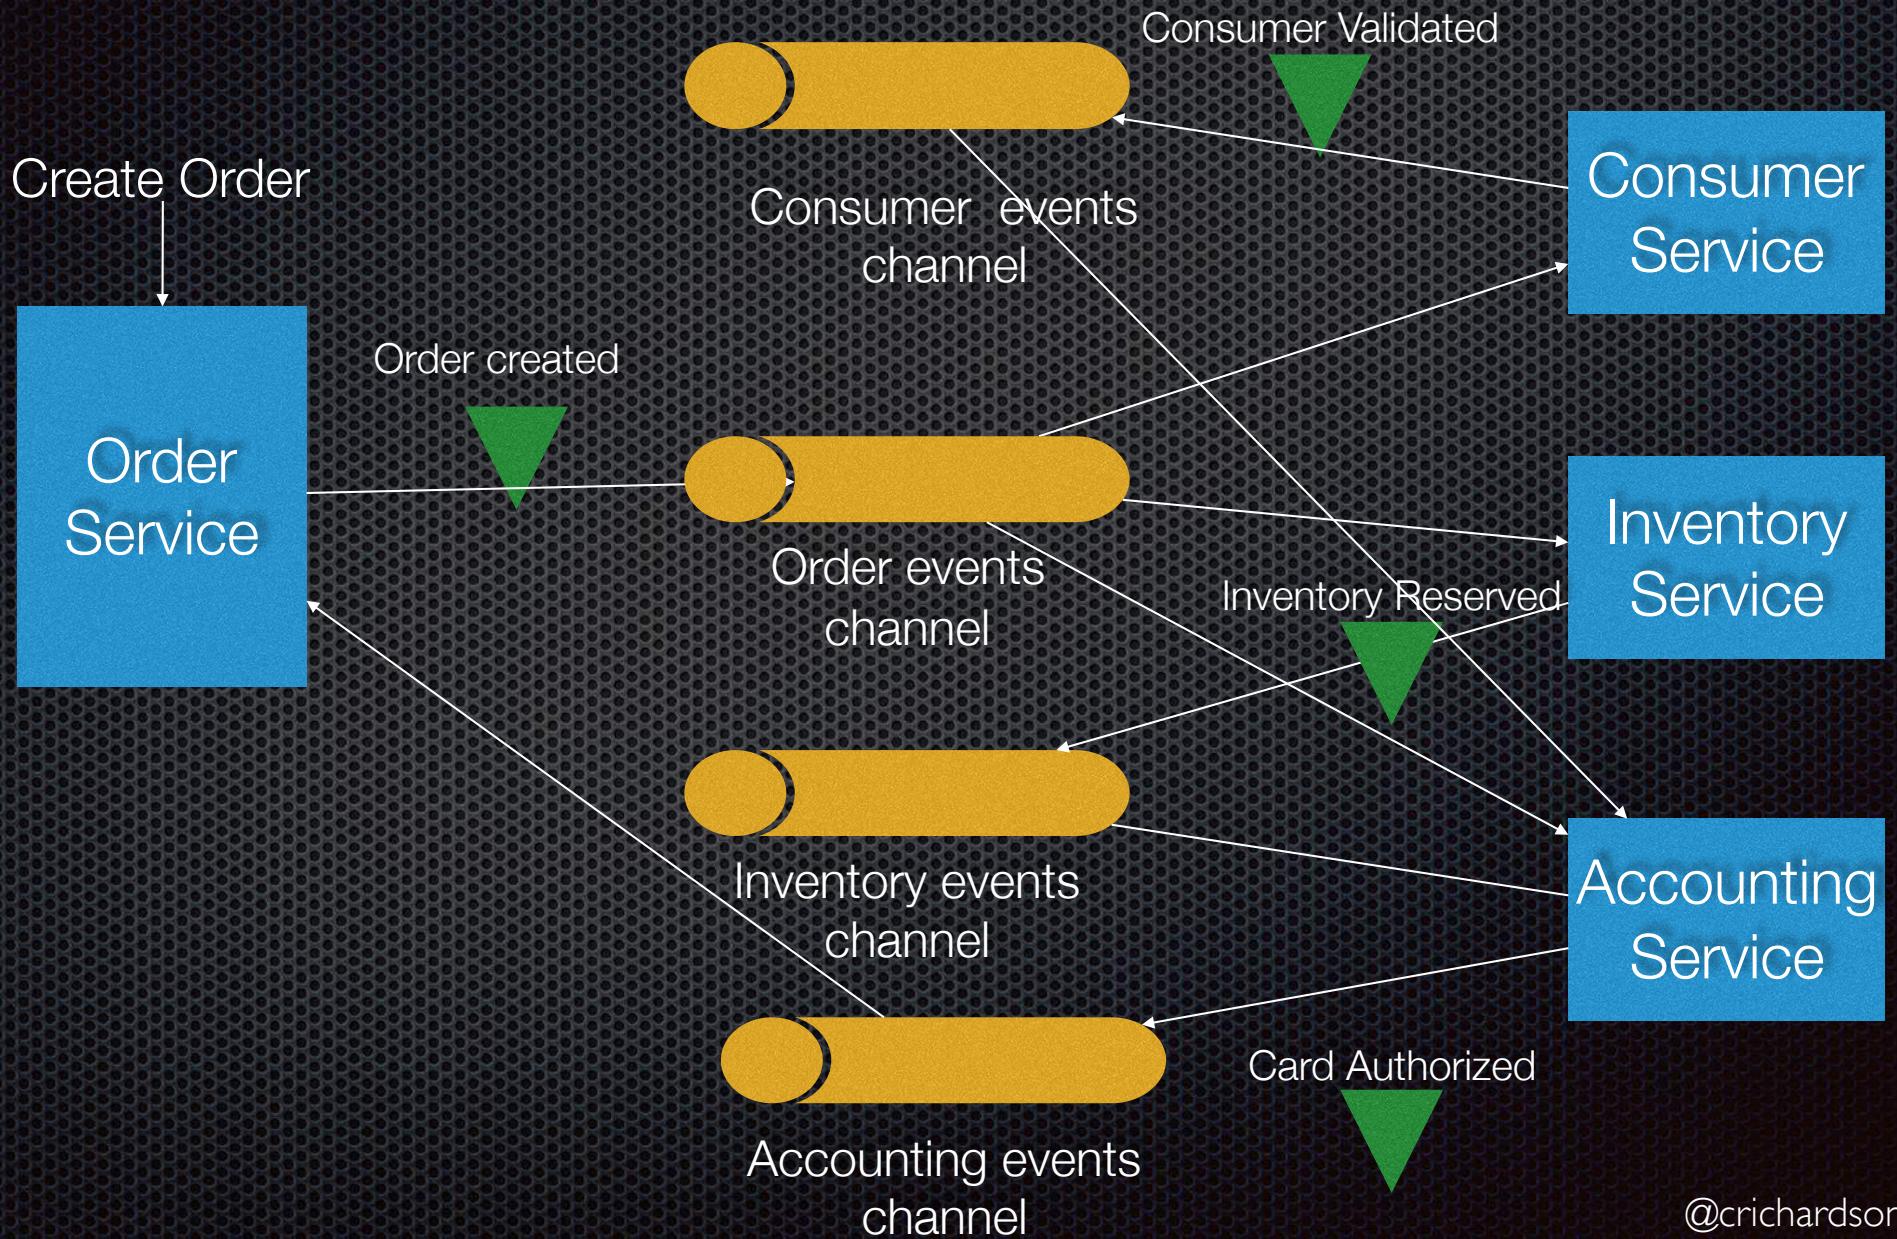


Attempt to Reserve credit

Publish event on success

Publish event on failure

More complex choreography example



Benefits and drawbacks of choreography

Benefits

- Simple, especially when using event sourcing
- Participants are loosely coupled

Drawbacks

- Cyclic dependencies - services listen to each other's events
- Overloads domain objects, e.g. Order and Customer **know** too much
- Events = indirect way to make something happen

Option #2: Orchestration-based saga coordination

createOrder()

CreateOrderSaga

Invokes

Order Service

Local transaction

createOrder()

Order

state=PENDING

Invokes

Customer Service

Local transaction

reserveCredit()

Customer

Invokes

Order Service

Local transaction

approve
order()

Order

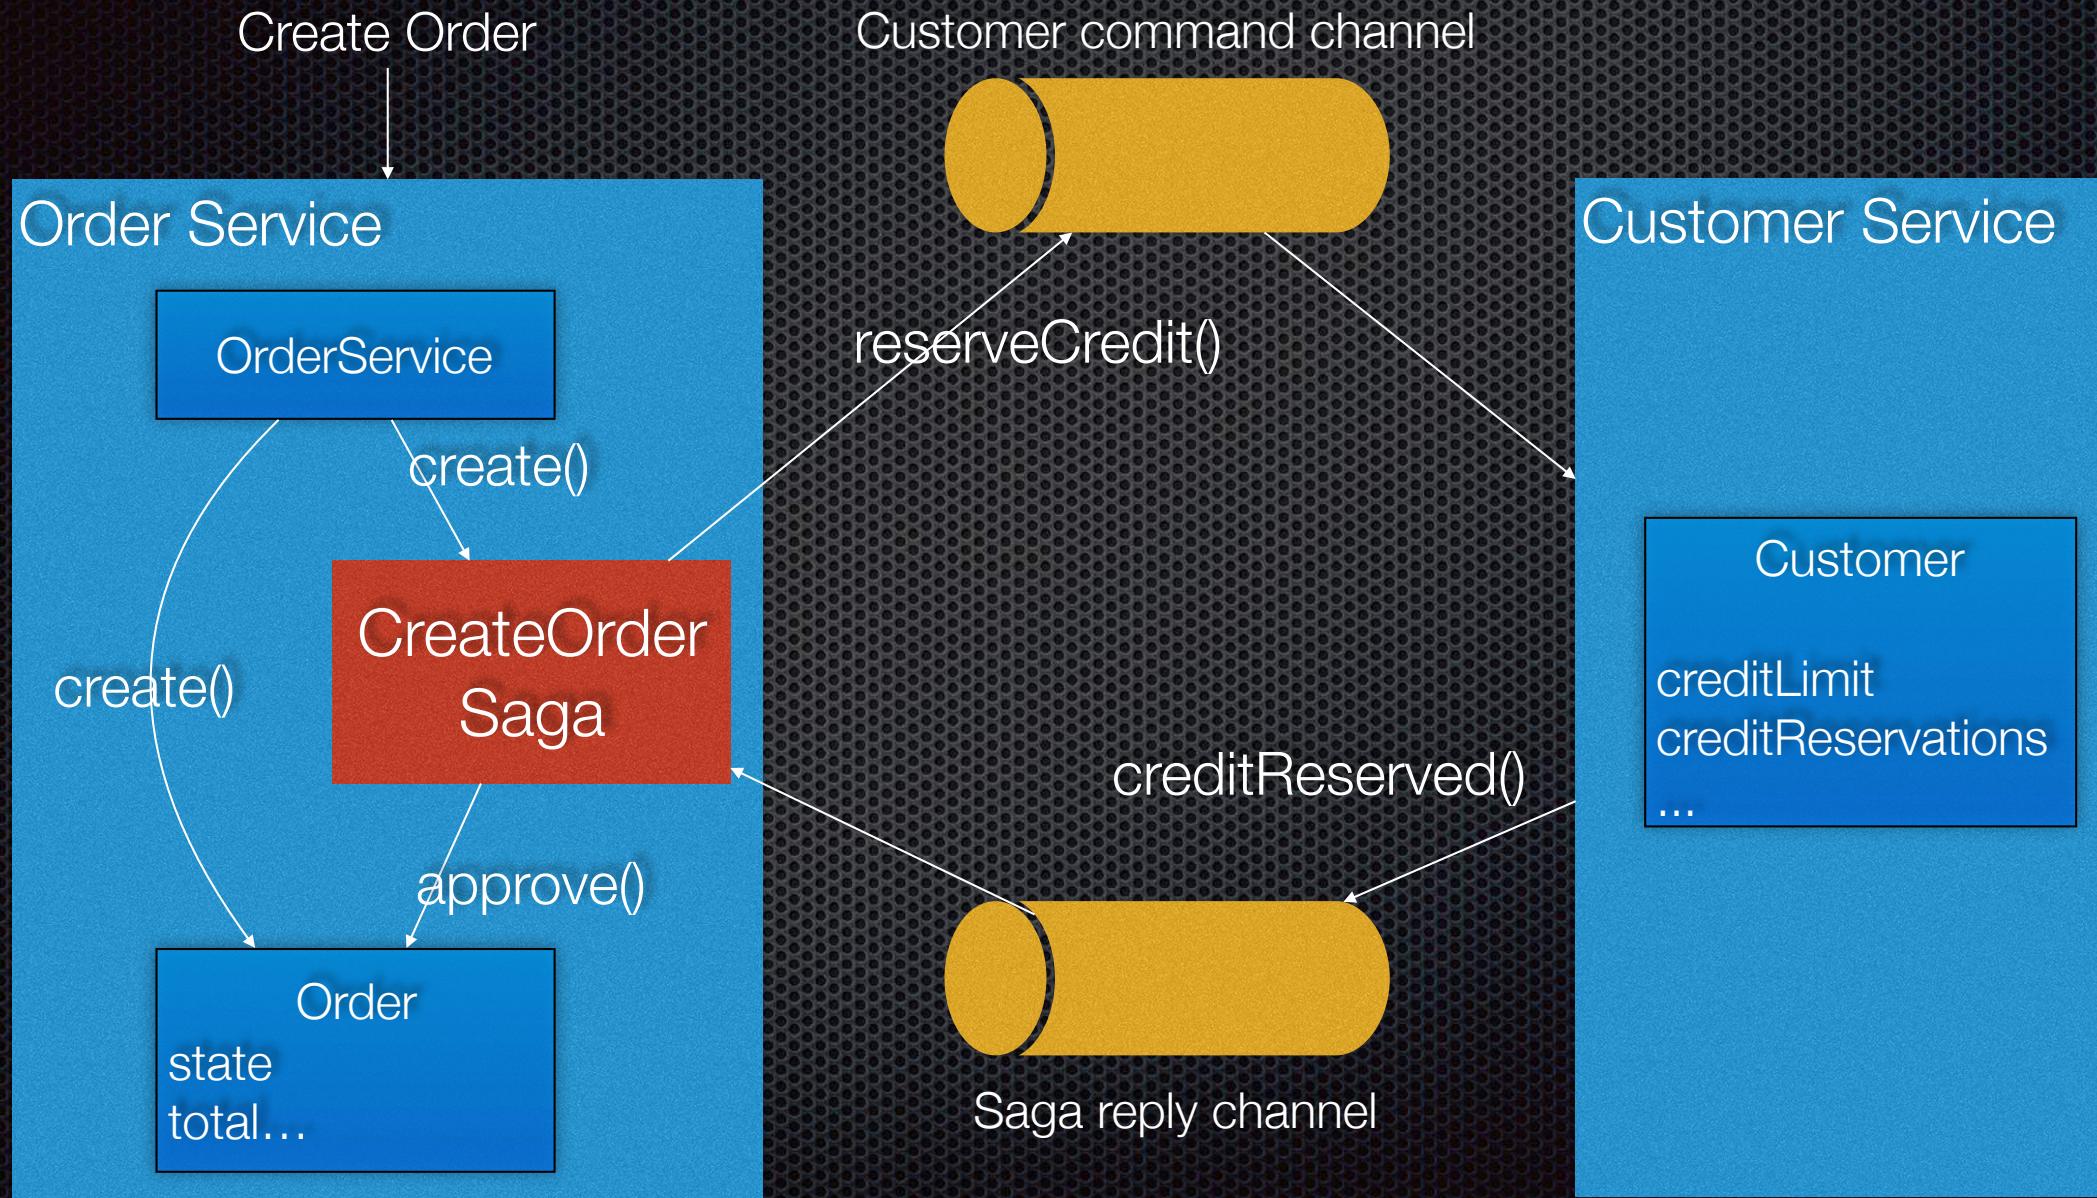
state=APPROVED

A saga (orchestrator)
is a **persistent object**
that
tracks the state of the saga
and
invokes the participants

Saga orchestrator behavior

- ❖ On create:
 - ❖ Invokes a saga participant
 - ❖ Persists state in database
 - ❖ Wait for a reply
- ❖ On reply:
 - ❖ Load state from database
 - ❖ Determine which saga participant to invoke next
 - ❖ Invokes saga participant
 - ❖ Updates its state
 - ❖ Persists updated state
 - ❖ Wait for a reply
 - ❖ ...

CreateOrderSaga orchestrator



CreateOrderSaga definition

```
public class CreateOrderSaga implements SimpleSaga<CreateOrderSagaData> {  
  
    private SagaDefinition<CreateOrderSagaData> sagaDefinition =  
        step()  
            .withCompensation(this::reject)  
        .step()  
            .invokeParticipant(this::reserveCredit)  
        .step()  
            .invokeParticipant(this::approve)  
        .build();  
  
    @Override  
    public SagaDefinition<CreateOrderSagaData> getSagaDefinition() { return this.sagaDefinition; }  
  
    private CommandWithDestination reserveCredit(CreateOrderSagaData data) {  
        long orderId = data.getOrderId();  
        Long customerId = data.getOrderDetails().getCustomerId();  
        Money orderTotal = data.getOrderDetails().getOrderTotal();  
        return send(new ReserveCreditCommand(customerId, orderId, orderTotal))  
            .to("customerService")  
            .build();  
    }  
}
```

Saga's Data

Sequence of steps

step = (T_i , C_i)

Build command to send

Customer Service command handler

```
public class CustomerCommandHandler {  
  
    @Autowired  
    private CustomerRepository customerRepository;  
  
    public CommandHandlers commandHandlerDefinitions() {  
        return SagaCommandHandlersBuilder  
            .fromChannel("customerService")  
            .onMessage(ReserveCreditCommand.class, this::reserveCredit)  
            .build();  
    }  
  
    public Message reserveCredit(CommandMessage<ReserveCreditCommand> cm) {  
        ReserveCreditCommand cmd = cm.getCommand();  
        long customerId = cmd.getCustomerId();  
        Customer customer = customerRepository.findOne(customerId);  
        try {  
            customer.reserveCredit(cmd.getOrderId(), cmd.getOrderTotal());  
            return withSuccess(new CustomerCreditReserved());  
        } catch (CustomerCreditLimitExceededException e) {  
            return withFailure(new CustomerCreditReservationFailed());  
        }  
    }  
}
```

Route command to handler

Make reply message

Reserve credit

@crichardson

Eventuate Tram Sagas

- Open-source Saga orchestration framework
- Currently for Java
- <https://github.com/eventuate-tram/eventuate-tram-sagas>
- <https://github.com/eventuate-tram/eventuate-tram-sagas-examples-customers-and-orders>

Benefits and drawbacks of orchestration

Benefits

- Centralized coordination logic is easier to understand
- Reduced coupling, e.g.
Customer knows less.
Simply has API
- Reduces cyclic dependencies

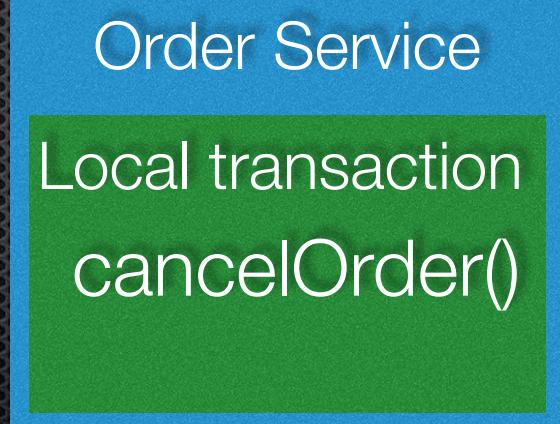
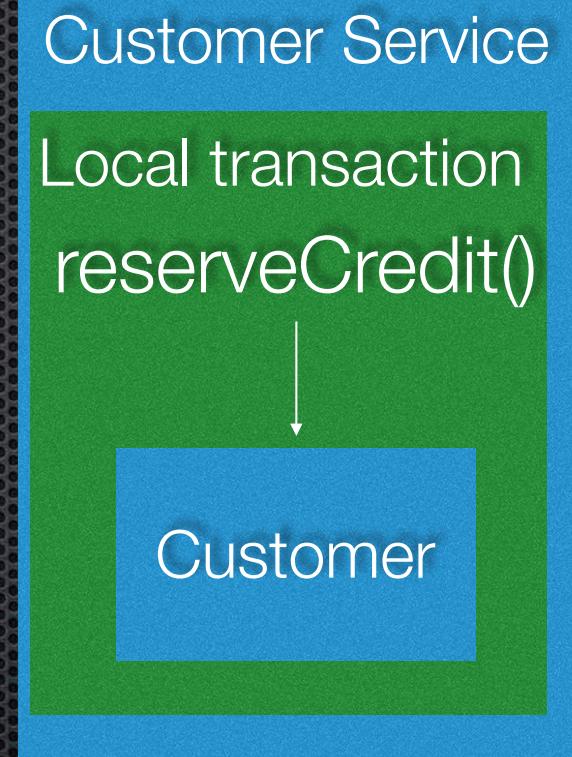
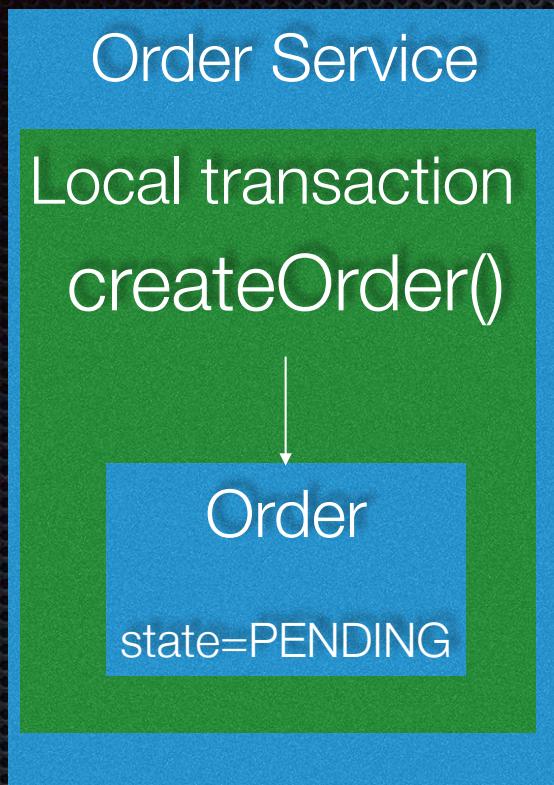
Drawbacks

- Risk of smart sagas directing dumb services

Agenda

- ACID is not an option
- Overview of sagas
- Coordinating sagas
- Countermeasures for data anomalies
- Using reliable and transactional messaging

Lack of isolation ⇒ complicates business logic



Time

How to cancel a PENDING Order?

- ☒ Don't ⇒ throw an OrderNotCancellableException
 - ☒ Questionable user experience
- ☒ "Interrupt" the Create Order saga?
 - ☒ Cancel Order Saga: set order.state = CANCELLED
 - ☒ Causes Create Order Saga to rollback
 - ☒ But is that enough to cancel the order?
- ☒ Cancel Order saga waits for the Create Order saga to complete?
 - ☒ Suspiciously like a distributed lock
 - ☒ But perhaps that is ok

Countermeasure Transaction Model

SOFTWARE—PRACTICE AND EXPERIENCE, VOL. 28(1), 77–98 (JANUARY 1998)

Semantic ACID Properties in Multidatabases Using Remote Procedure Calls and Update Propagations

lars frank¹ and torben u. zahle²

<http://bit.ly/semantic-acid-ctm> - paywall 😞

Sagas are ACD

- **Atomicity**
 - Saga implementation ensures that all transactions are executed **OR** all are compensated
- **Consistency**
 - Referential integrity within a service handled by local databases
 - Referential integrity across services handled by application
- **Durability**
 - Durability handled by local databases

Lack of $I \rightarrow$ anomalies

Outcome of
concurrent execution
!=
a sequential execution

Sounds scary
BUT
It's common to relax isolation
to improve performance

Anomaly: Lost update

Create
Order
Saga

Ti: Create
Order

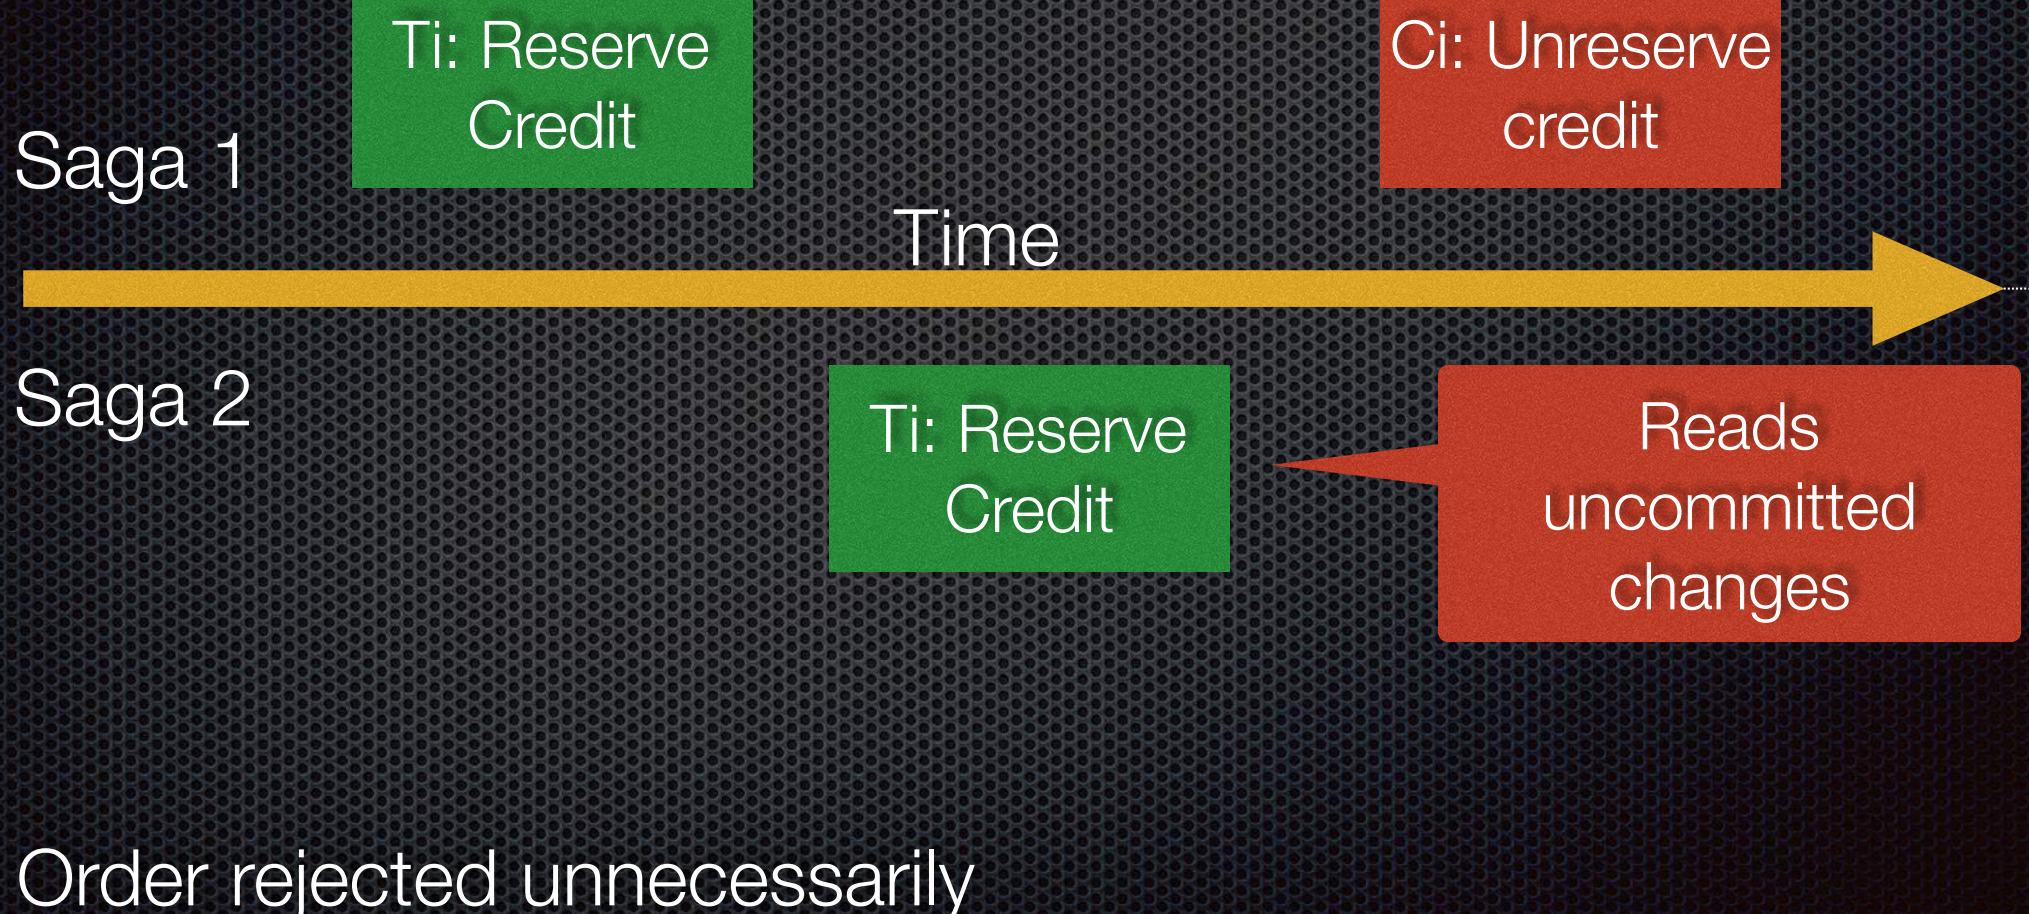
Cancel
Order
Saga

Time

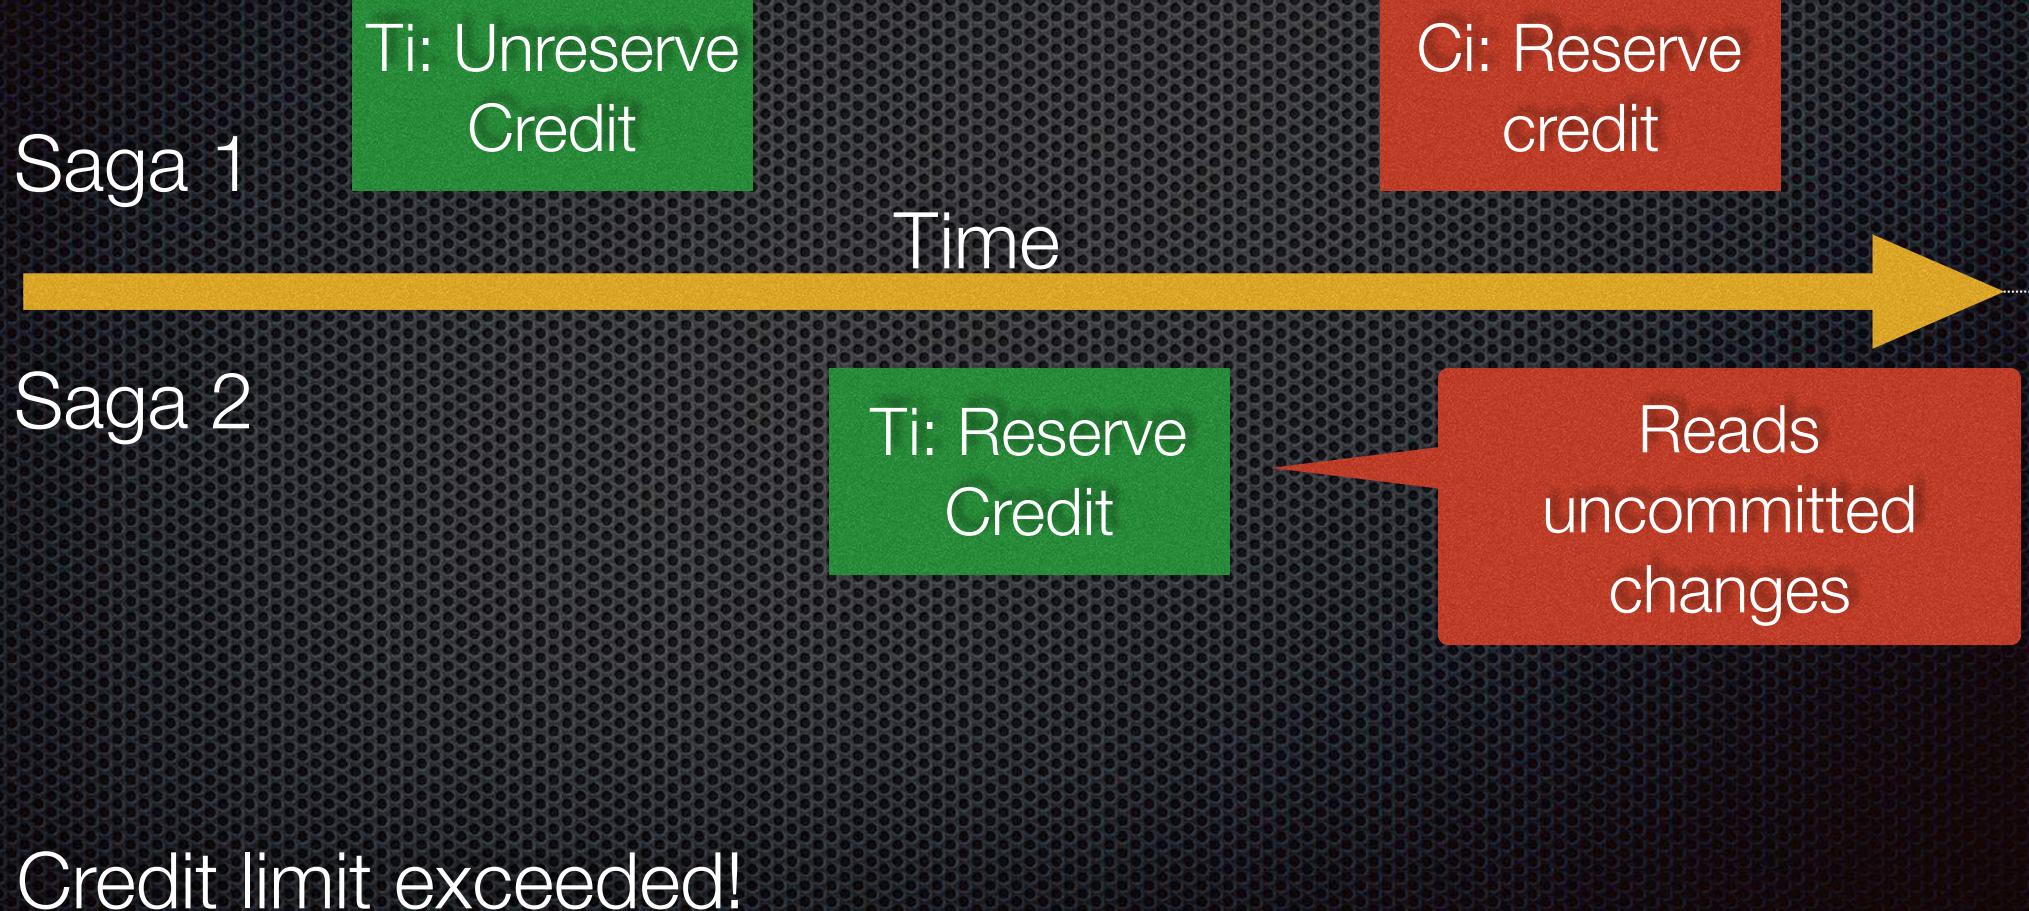
Tj: Approve
Order

Overwrites
cancelled order

Anomaly: Dirty reads . . .



... Anomaly: Dirty reads



Anomaly: non-repeatable/ fuzzy read

Saga 1

Ti: Begin
Revise Order

Tj: Finalize
Revise Order

Saga 2

Ti: *Update*
Order

Order has changed
since Ti

Time



Countermeasures for **reducing^{*}** impact of isolation anomalies...

i.e. good enough, eventually consistent application

Saga structure

Compensatable transactions



T1

C1

T2

C2

...

T_{n+1}

Pivot transaction = GO/NO GO



T_{n+2}

Retriable transactions that can't fail



...

Countermeasure: Semantic lock

- Compensatable transaction sets flag, retriable transaction releases it
- Indicates a possible dirty read
- Flag = lock:
 - prevents other transactions from accessing it
 - Require deadlock detection, e.g. timeout
- Flag = warning - treat the data differently, e.g.
 - Order.state = PENDING
 - a pending deposit

Order.state =
PENDING

Order.state =
REJECTED

Create Pending
Order

Reject Order

...

...

...

Approve Order

Order.state =
APPROVED

Countermeasure: Commutative updates

- ▣ Commutative: $g(f(x)) = f(g(x))$
- ▣ For example:
 - ▣ Account.debit() compensates for Account.credit()
 - ▣ Account.credit() compensates for Account.debit()
- ▣ Avoids lost updates

Countermeasure: Pessimistic view

...

Increase avail.
credit

Cancel Order Delivery

Cancel Order

...

Reduce avail.
credit

Reorder saga
to
reduce **risk**

...

Cancel Order Delivery

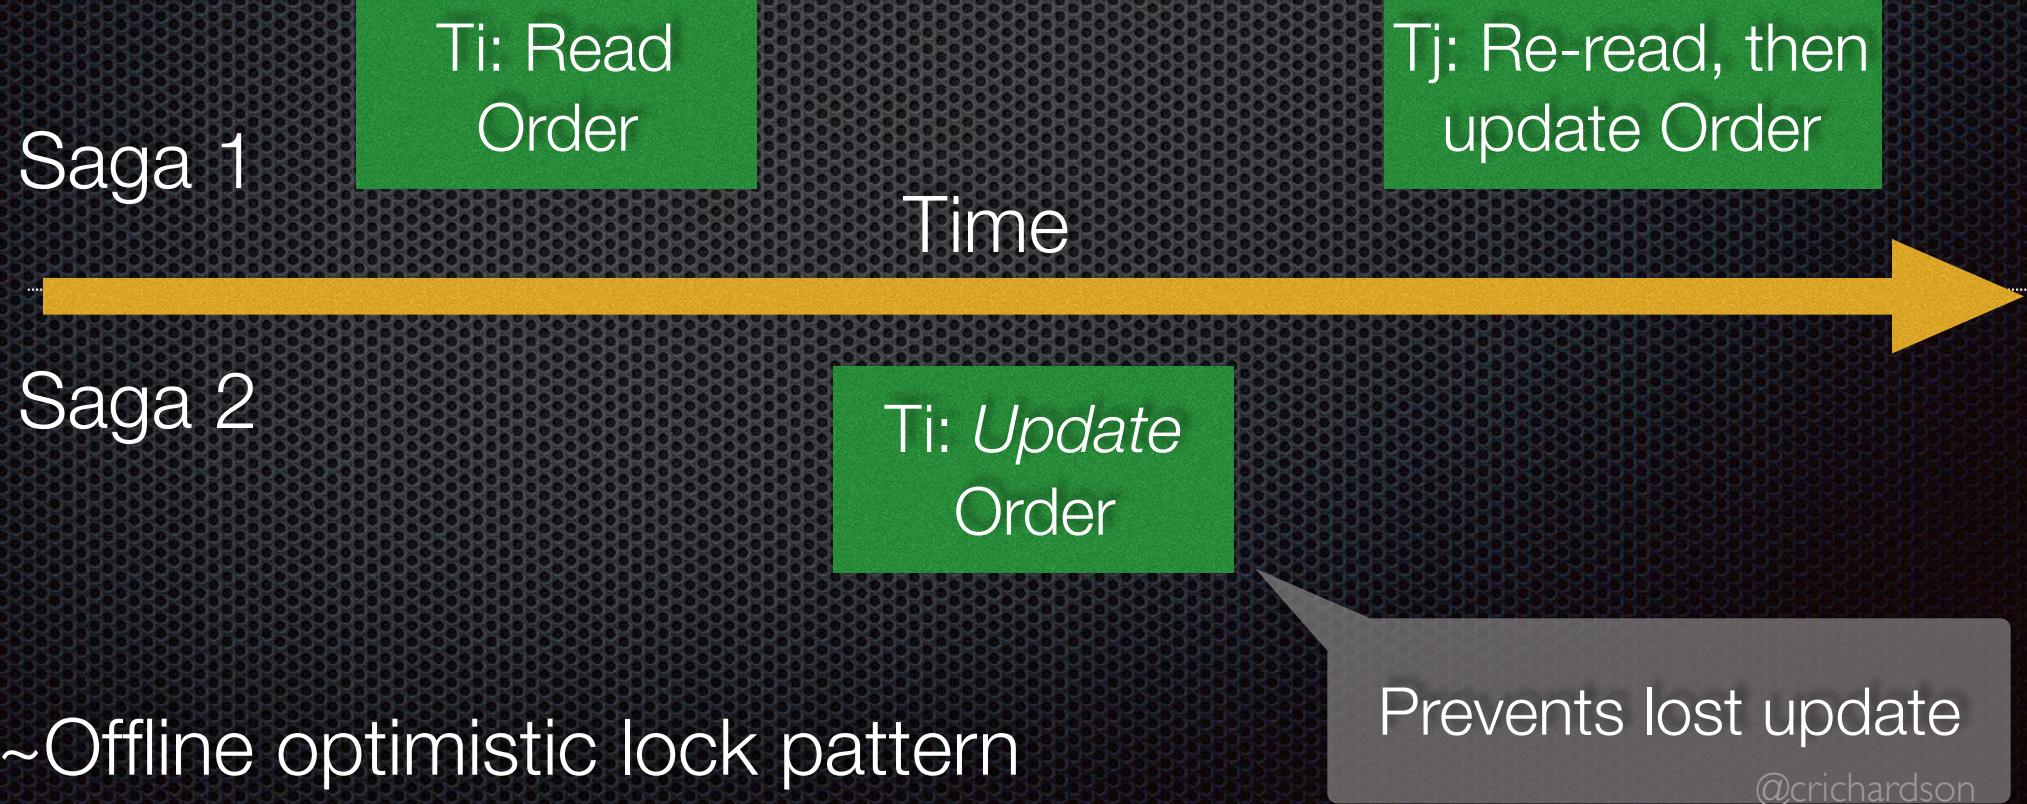
Cancel Order

Increase avail. credit

Won't be compensated,
so no dirty read

Countermeasure: Re-read value

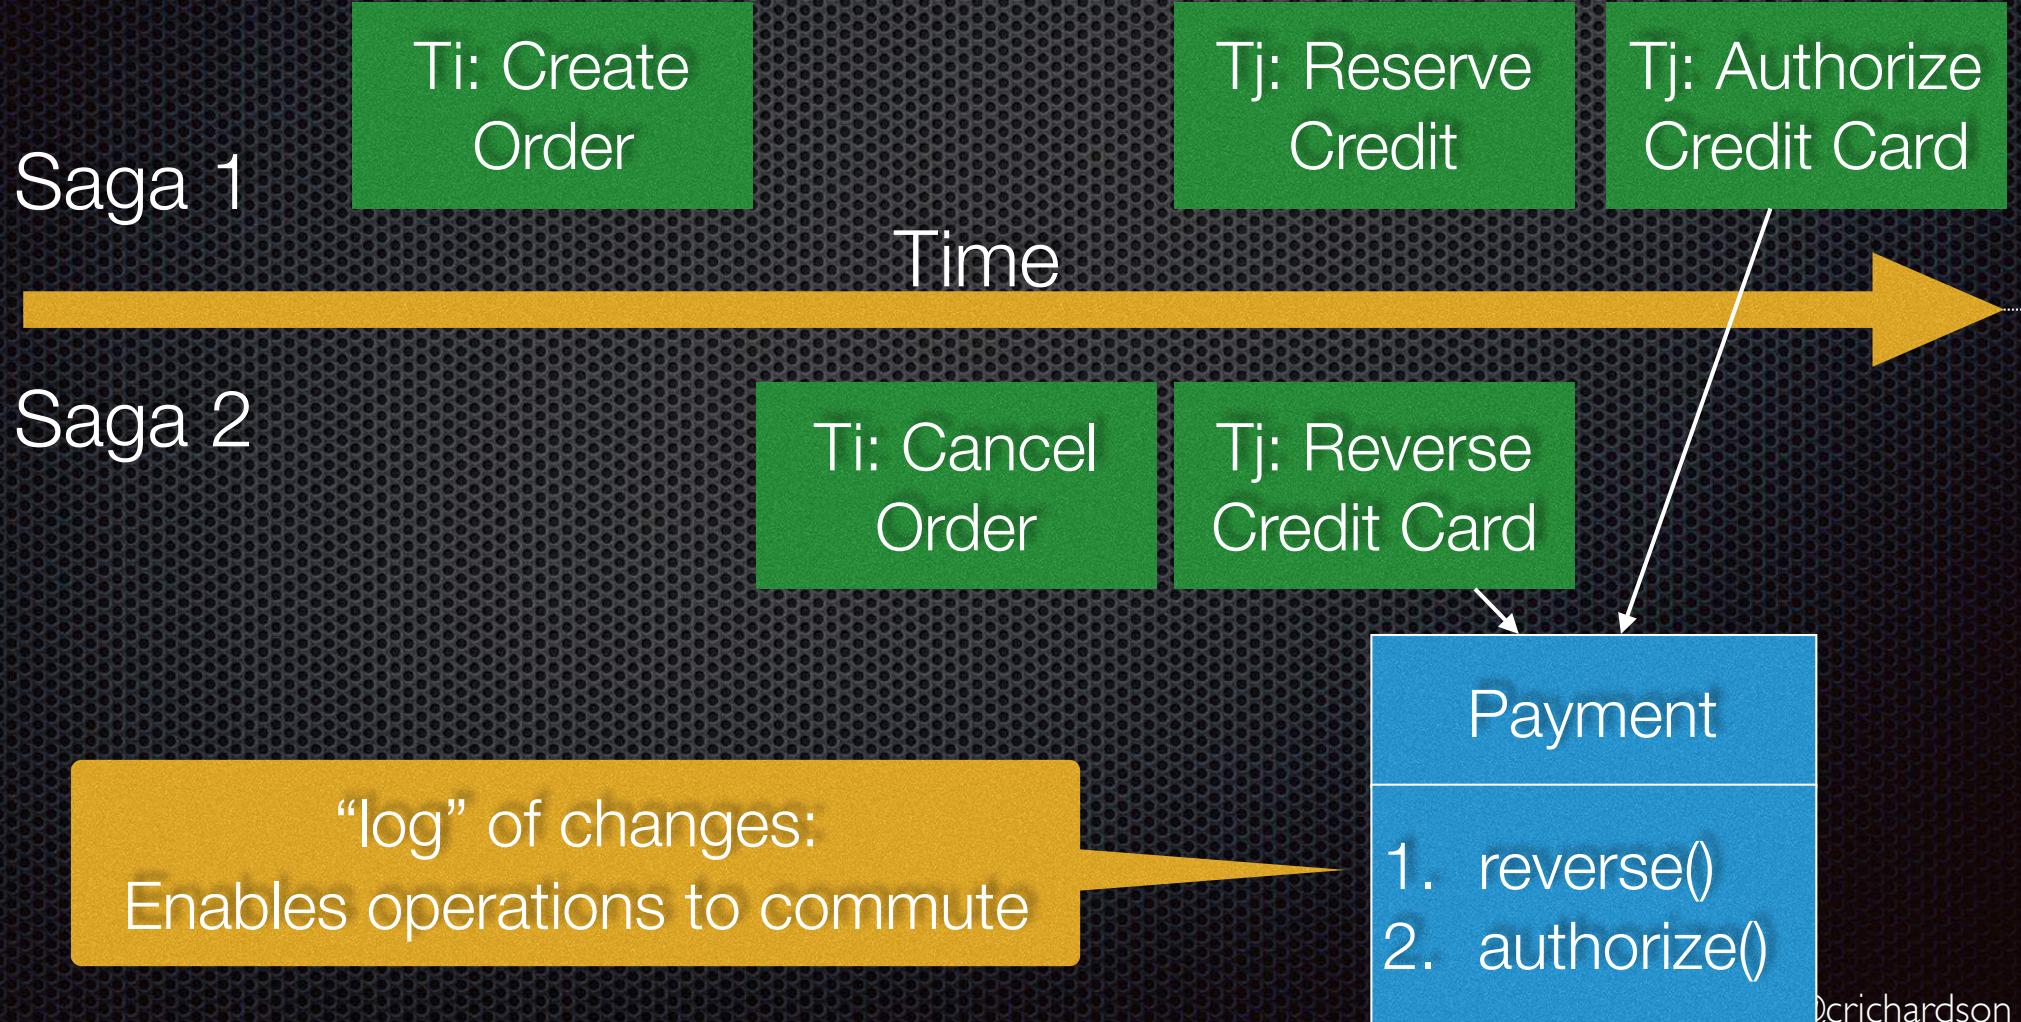
Verify unchanged, possibly restart



~Offline optimistic lock pattern

Prevents lost update
@crichtson

Countermeasure: version file



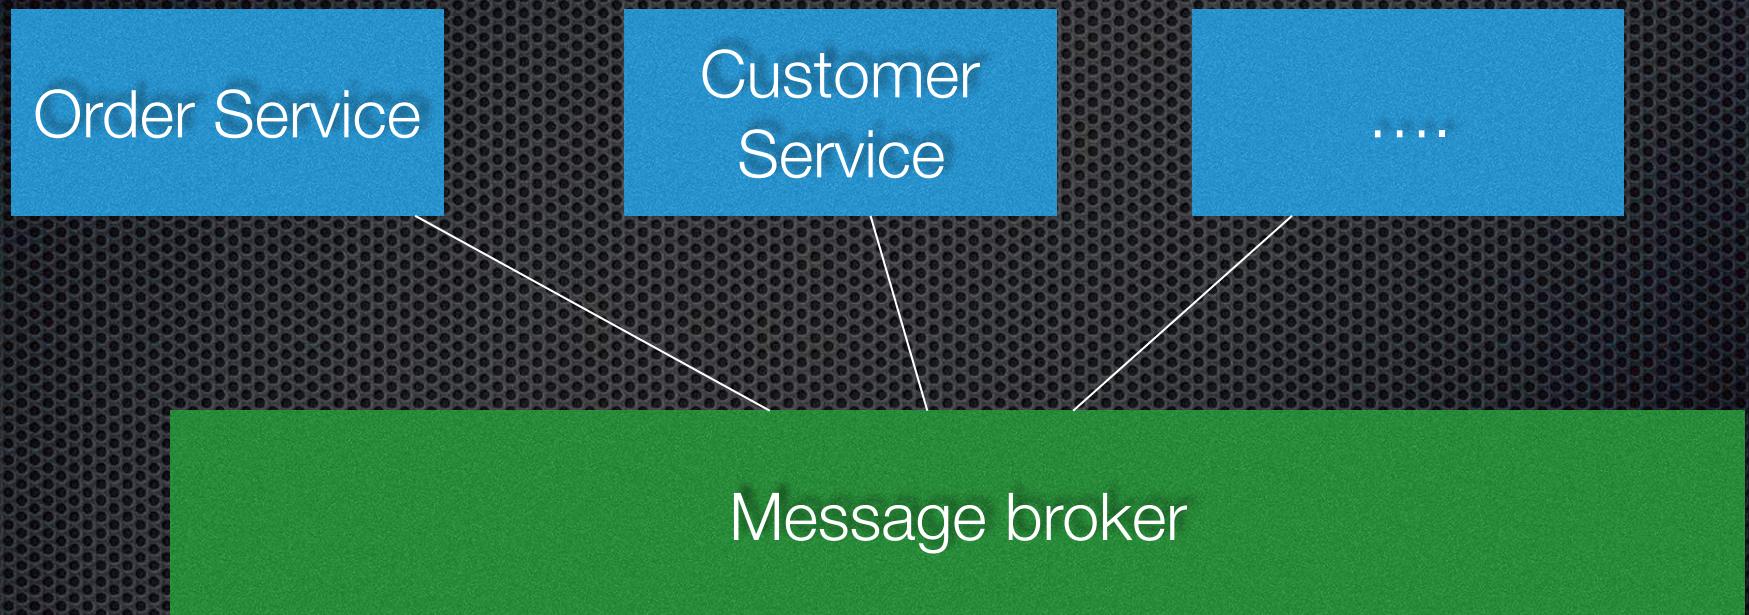
Countermeasure: By value

- ❖ Business risk determines strategy
- ❖ Low risk => semantic ACID
- ❖ High risk => use 2PC/distributed transaction

Agenda

- ACID is not an option
- Overview of sagas
- Coordinating sagas
- Countermeasures for data anomalies
- Using reliable and transactional messaging

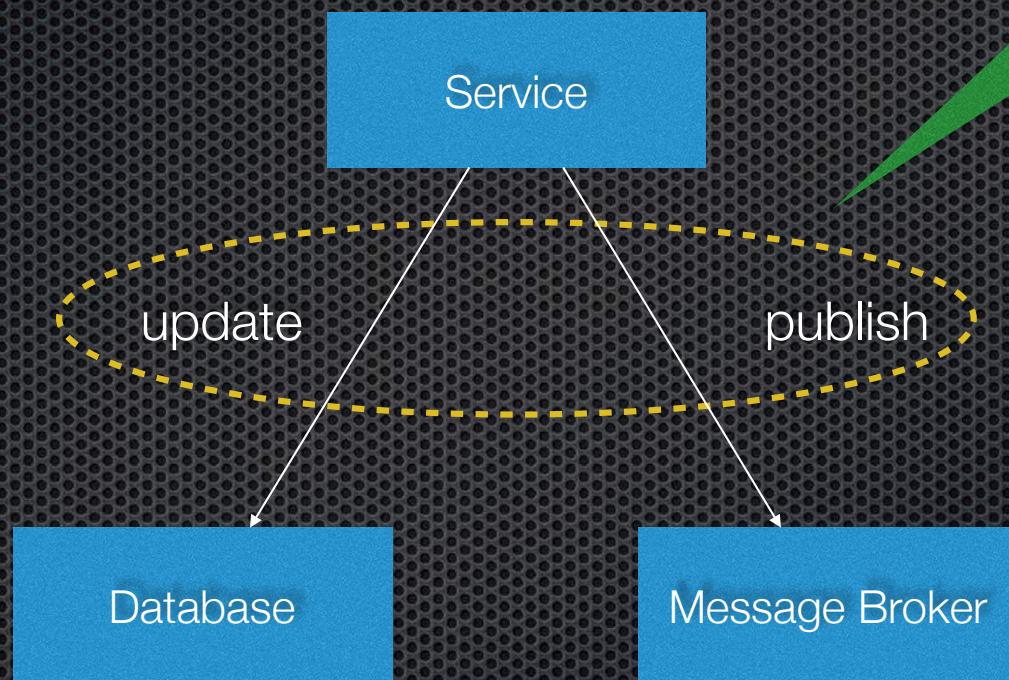
Use asynchronous, broker-based messaging



Guaranteed delivery ensures a saga complete when its participants are temporarily unavailable

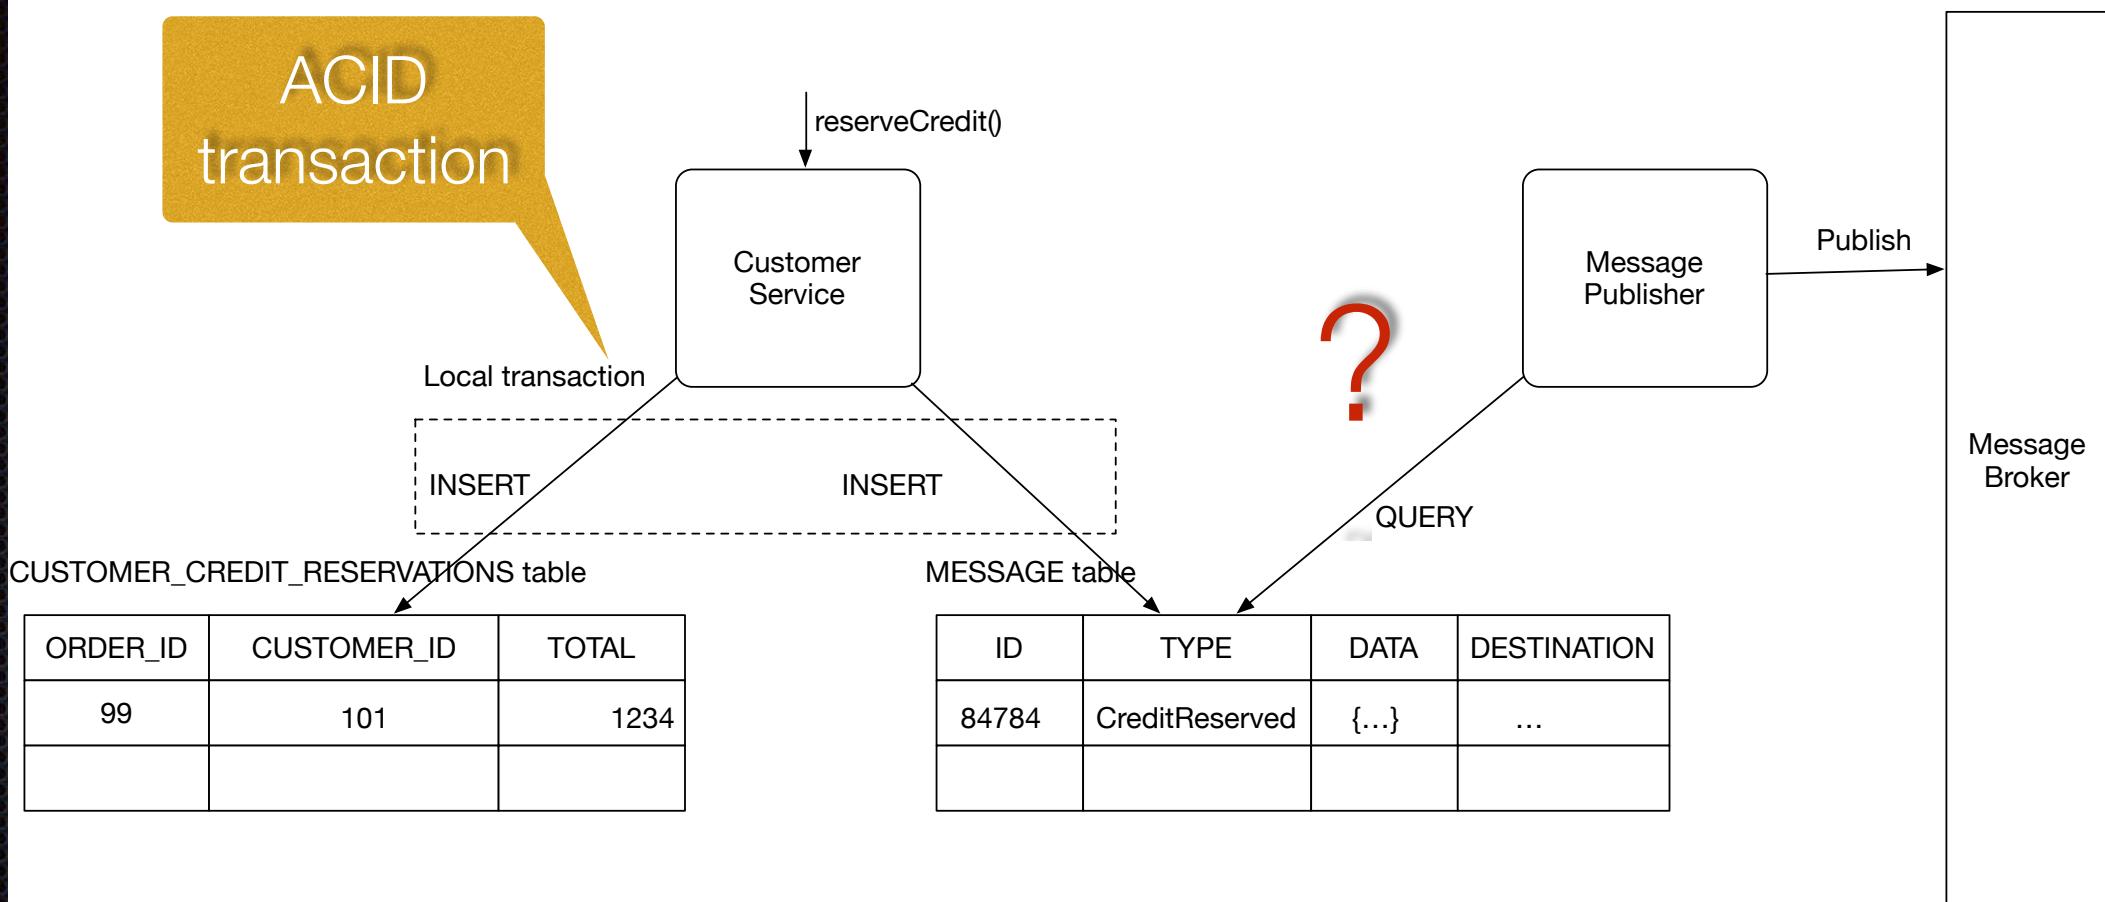
Messaging must be transactional

How to
make **atomic**
without 2PC?



Option #1: Use database table as a message queue

Option #1: Use database table as a message queue



- See BASE: An Acid Alternative, <http://bit.ly/ebaybase>

Publishing messages by polling the MESSAGE table

```
SELECT *  
FROM MESSAGES  
WHERE ...
```

Message
Publisher

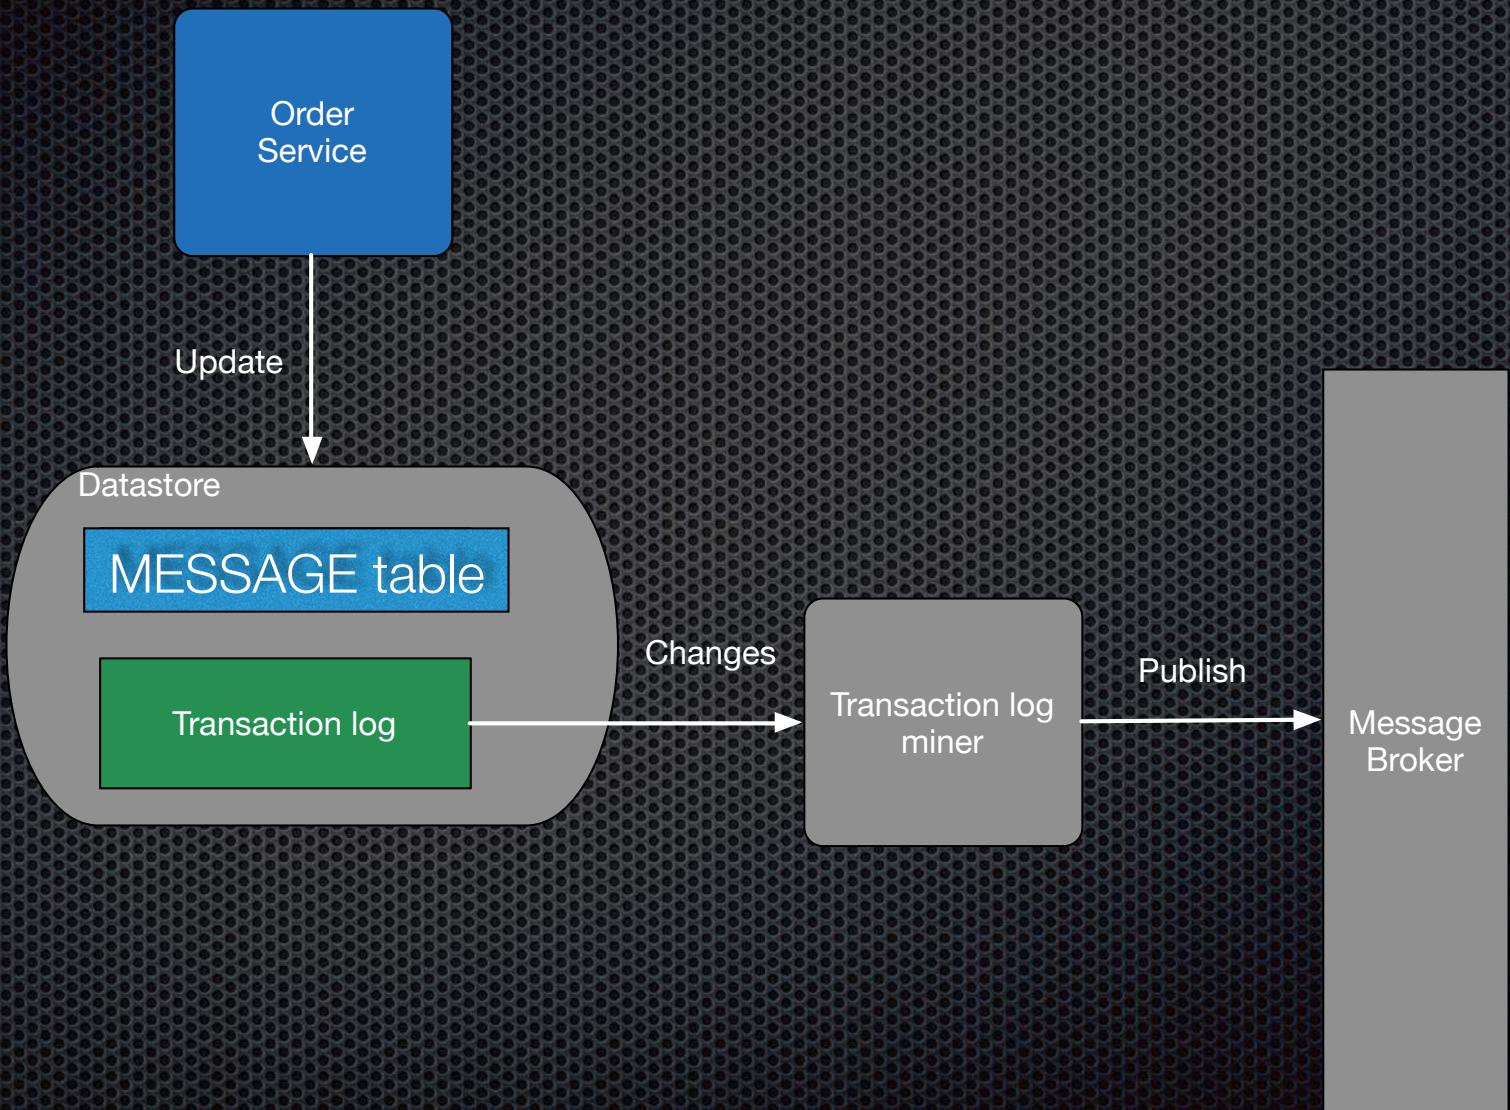
```
DELETE ...
```

Publish message

Message table

Message table

Transaction log tailing



Transaction log tailing

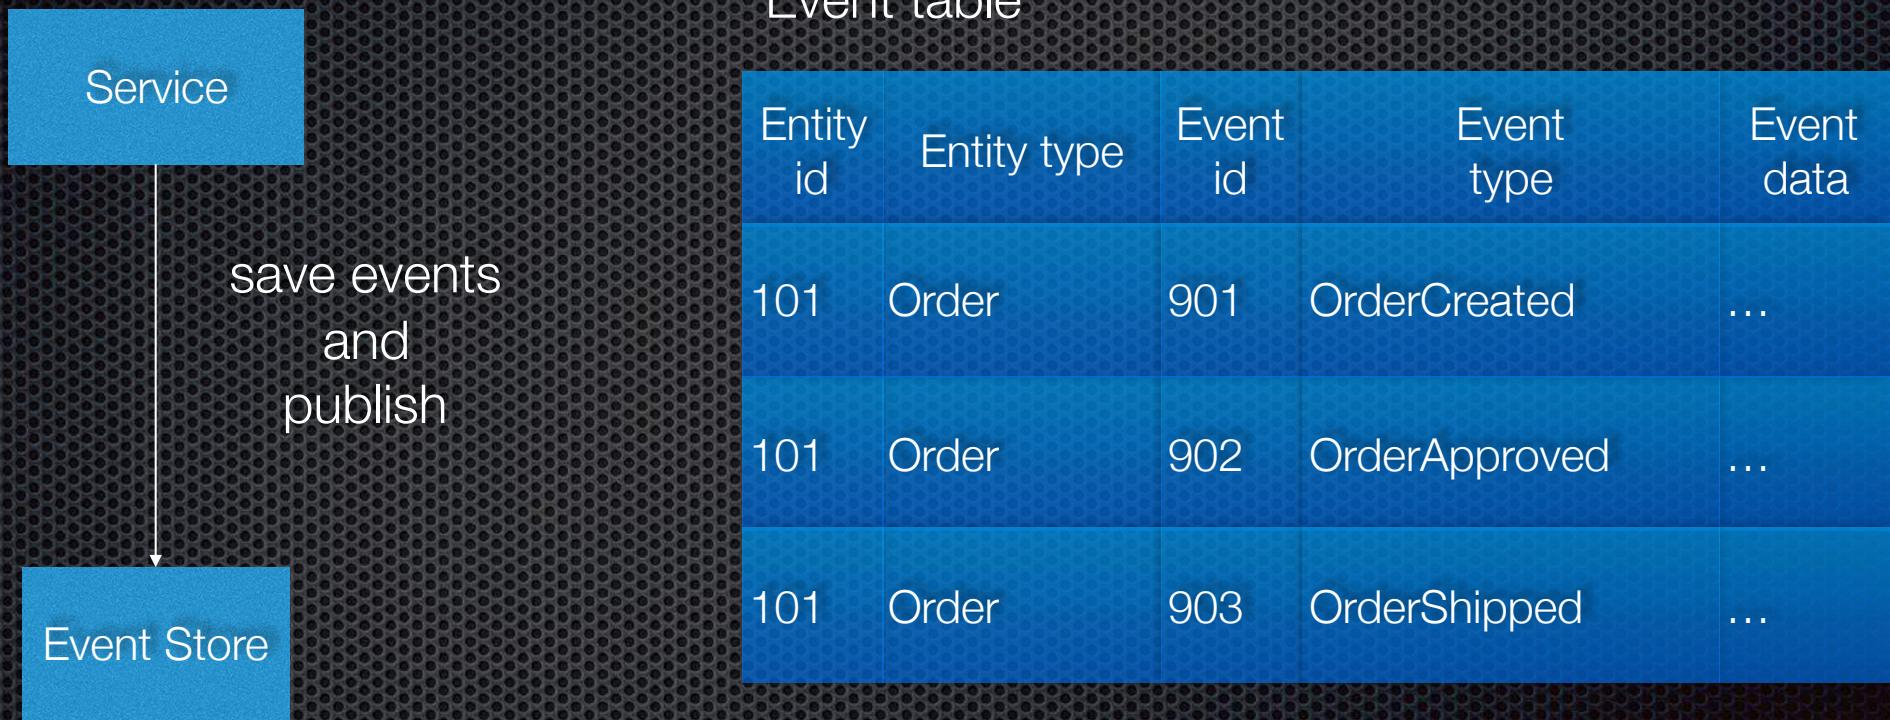
- Eventuate Local - reads MySQL binlog
- Oracle Golden Gate
- AWS DynamoDB streams
- MongoDB - Read the oplog

Transaction log tailing: benefits and drawbacks

- **Benefits**
 - No 2PC
 - No application changes required
 - Guaranteed to be accurate
- **Drawbacks**
 - Obscure
 - Database specific solutions
 - Tricky to avoid duplicate publishing

Option #2: Event sourcing: event-centric persistence

Event sourcing: persists an object as a sequence of events



Every state change ⇒ event

Guarantees:

state change

->

event is published

Implementing choreography-based sagas is straightforward

Preserves history of domain objects

Supports temporal queries

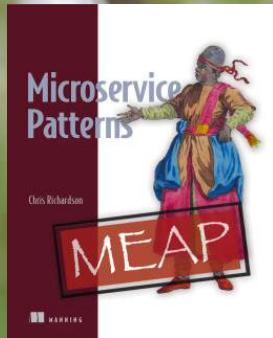
Built-in auditing

Summary

- Microservices tackle complexity and accelerate development
- Database per service is essential for loose coupling
- Use ACID transactions **within** services
- Use orchestration-based or choreography-based sagas to maintain data consistency **across** services
- Use countermeasures to reduce impact of anomalies caused by lack of isolation



@crichtardson chris@chrisrichardson.net



40%
discount
with code
ctwsaturn18

Questions?

<http://learn.microservices.io>